

A whole bushel of AI things

<https://bit.ly/soma-econ>

Jonathan Soma

Knight Chair in Data Journalism

Columbia University Graduate School of Journalism

jonathansoma.com :: @dangerscarf :: js4571@columbia.edu

The Lede Program

An intensive, remote-friendly summer program from Columbia's Graduate School of Journalism designed to equip journalists and storytellers of all kinds with the skills needed to perform in-depth investigations

The Lede Program

Columbia Journalism School's Lede Program is a 10-week intensive on coding, data analysis, and visual storytelling. For 2025, the course will be once again operate with **in-person and remote options**.

No Prerequisites

We start from zero, **with no prior experience in data or coding necessary** – You'll learn it all along the way. If you can turn on your computer, you're ready to go!

10 weeks!

Not just for journalists

Pulling insights from data **isn't just for journalists** — our tools and skills are equally valuable to social scientists, policymakers, and academics of all stripes.

Real-World Skills

Tired of programming tutorials that are all about complex ways of putting a dozen numbers in order? We teach the **practical, real-world skills** in data analysis and visualization and sidestep the theory.

Have questions about Lede?

12 months!

COLUMBIA JOURNALISM SCHOOL

Academics ▾ People ▾ Professional Learning ▾ Centers ▾ Community ▾

Apply **Search**

M.S. Data Journalism

Journalism in the 21st century involves finding, collecting and analyzing data for storytelling, presentation and investigative reporting. The journalism school offers a Master of Science in Data Journalism for students interested in advanced skills.

APPLY

WHO SHOULD APPLY **CURRICULUM** **CLASSES**

The M.S. in Data Journalism provides the hands-on training needed to tell deeply reported data-driven stories in the public interest.

The current era needs journalists who can extract stories and meaning from data and massive information flows. This program trains students to be confident about using data in furtherance of the journalistic mission.

Applicants do not need to have experience with data or computation to enroll in this three-semester program. All students are required to attend foundational courses that allow those with no data experience to hone their skills in data acquisition, extraction and analysis.



Text inviting
friend to wedding



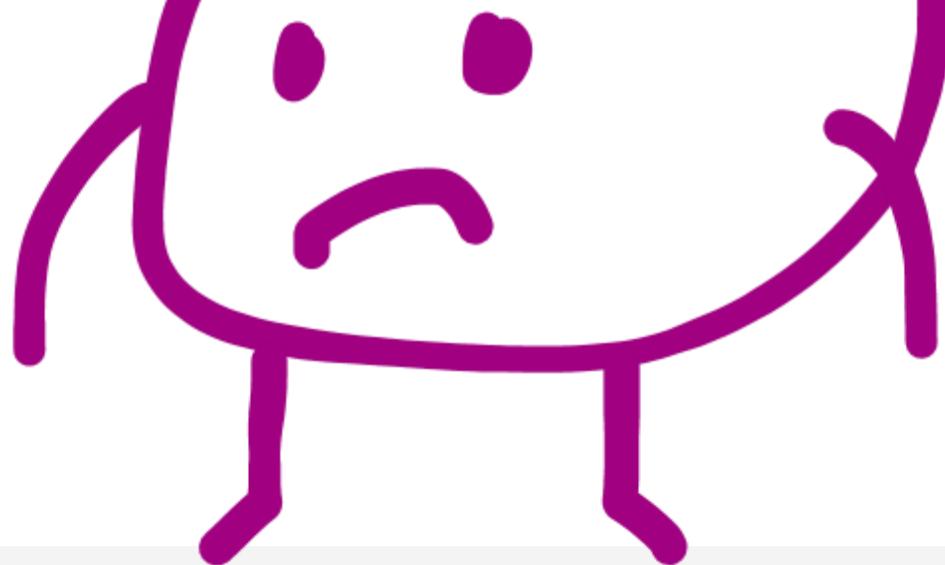
plain
superconductors



Experience
Seoul like a local



Design a fun
coding game



Message ChatGPT



ChatGPT can make mistakes. Check important info.



💬
Text inviting
friends to a wedding

💡
Explain
superconductors

💡
Experience
Seoul like a local

💡
Design a fun
coding game

✉️ Message ChatGPT



ChatGPT can make mistakes. Check important info.

The plan

1. A little bit of background on AI
2. Sorting documents
3. Using LLMs in spreadsheets
4. Creating structured data from messy text
5. Analyzing images and videos

youtube.com/playlist?list=PLewNEVDy7gq1_GPUaL0Q31QsiHP5ncAQ

YouTube BE

Search

Home

Shorts

Subscriptions

You >

Your channel

History

Playlists

Your videos

Watch later

Liked videos

Your clips

Dearest friend, I daresay I have not partaken of food in ages. I'm positively famished

Practical AI for Investigative Journalism

Jonathan Soma

Playlist • Public • 6 videos • 669 views

A six session series held in April 2024 about real-life use cases for journalism in (mostly investigative) jour...more

▶ Play all

Sort

Dearest friend, I daresay I have not partaken of food in ages. I'm positively famished 2:09:33

Sorting documents (Practical AI for Investigative Journalism, Session 1)

Jonathan Soma • 1.1K views • Streamed 2 months ago

Large language models don't understand facts or concepts, they only know statistical probability 2:14:32

Structured, validated data from LLMs (Practical AI for Investigative Journalism, Session 2)

Jonathan Soma • 743 views • Streamed 2 months ago

Evaluating Verifiability in Generative Search Engines 1:45:26

Why generative AI is a dead end for responsible journalism (Practical AI for Journalism, Session 3)

Jonathan Soma • 581 views • Streamed 2 months ago

AI, Hugging Face and non-chatbot models (Practical AI for Journalism, Session 4)

Jonathan Soma • 836 views • Streamed 2 months ago

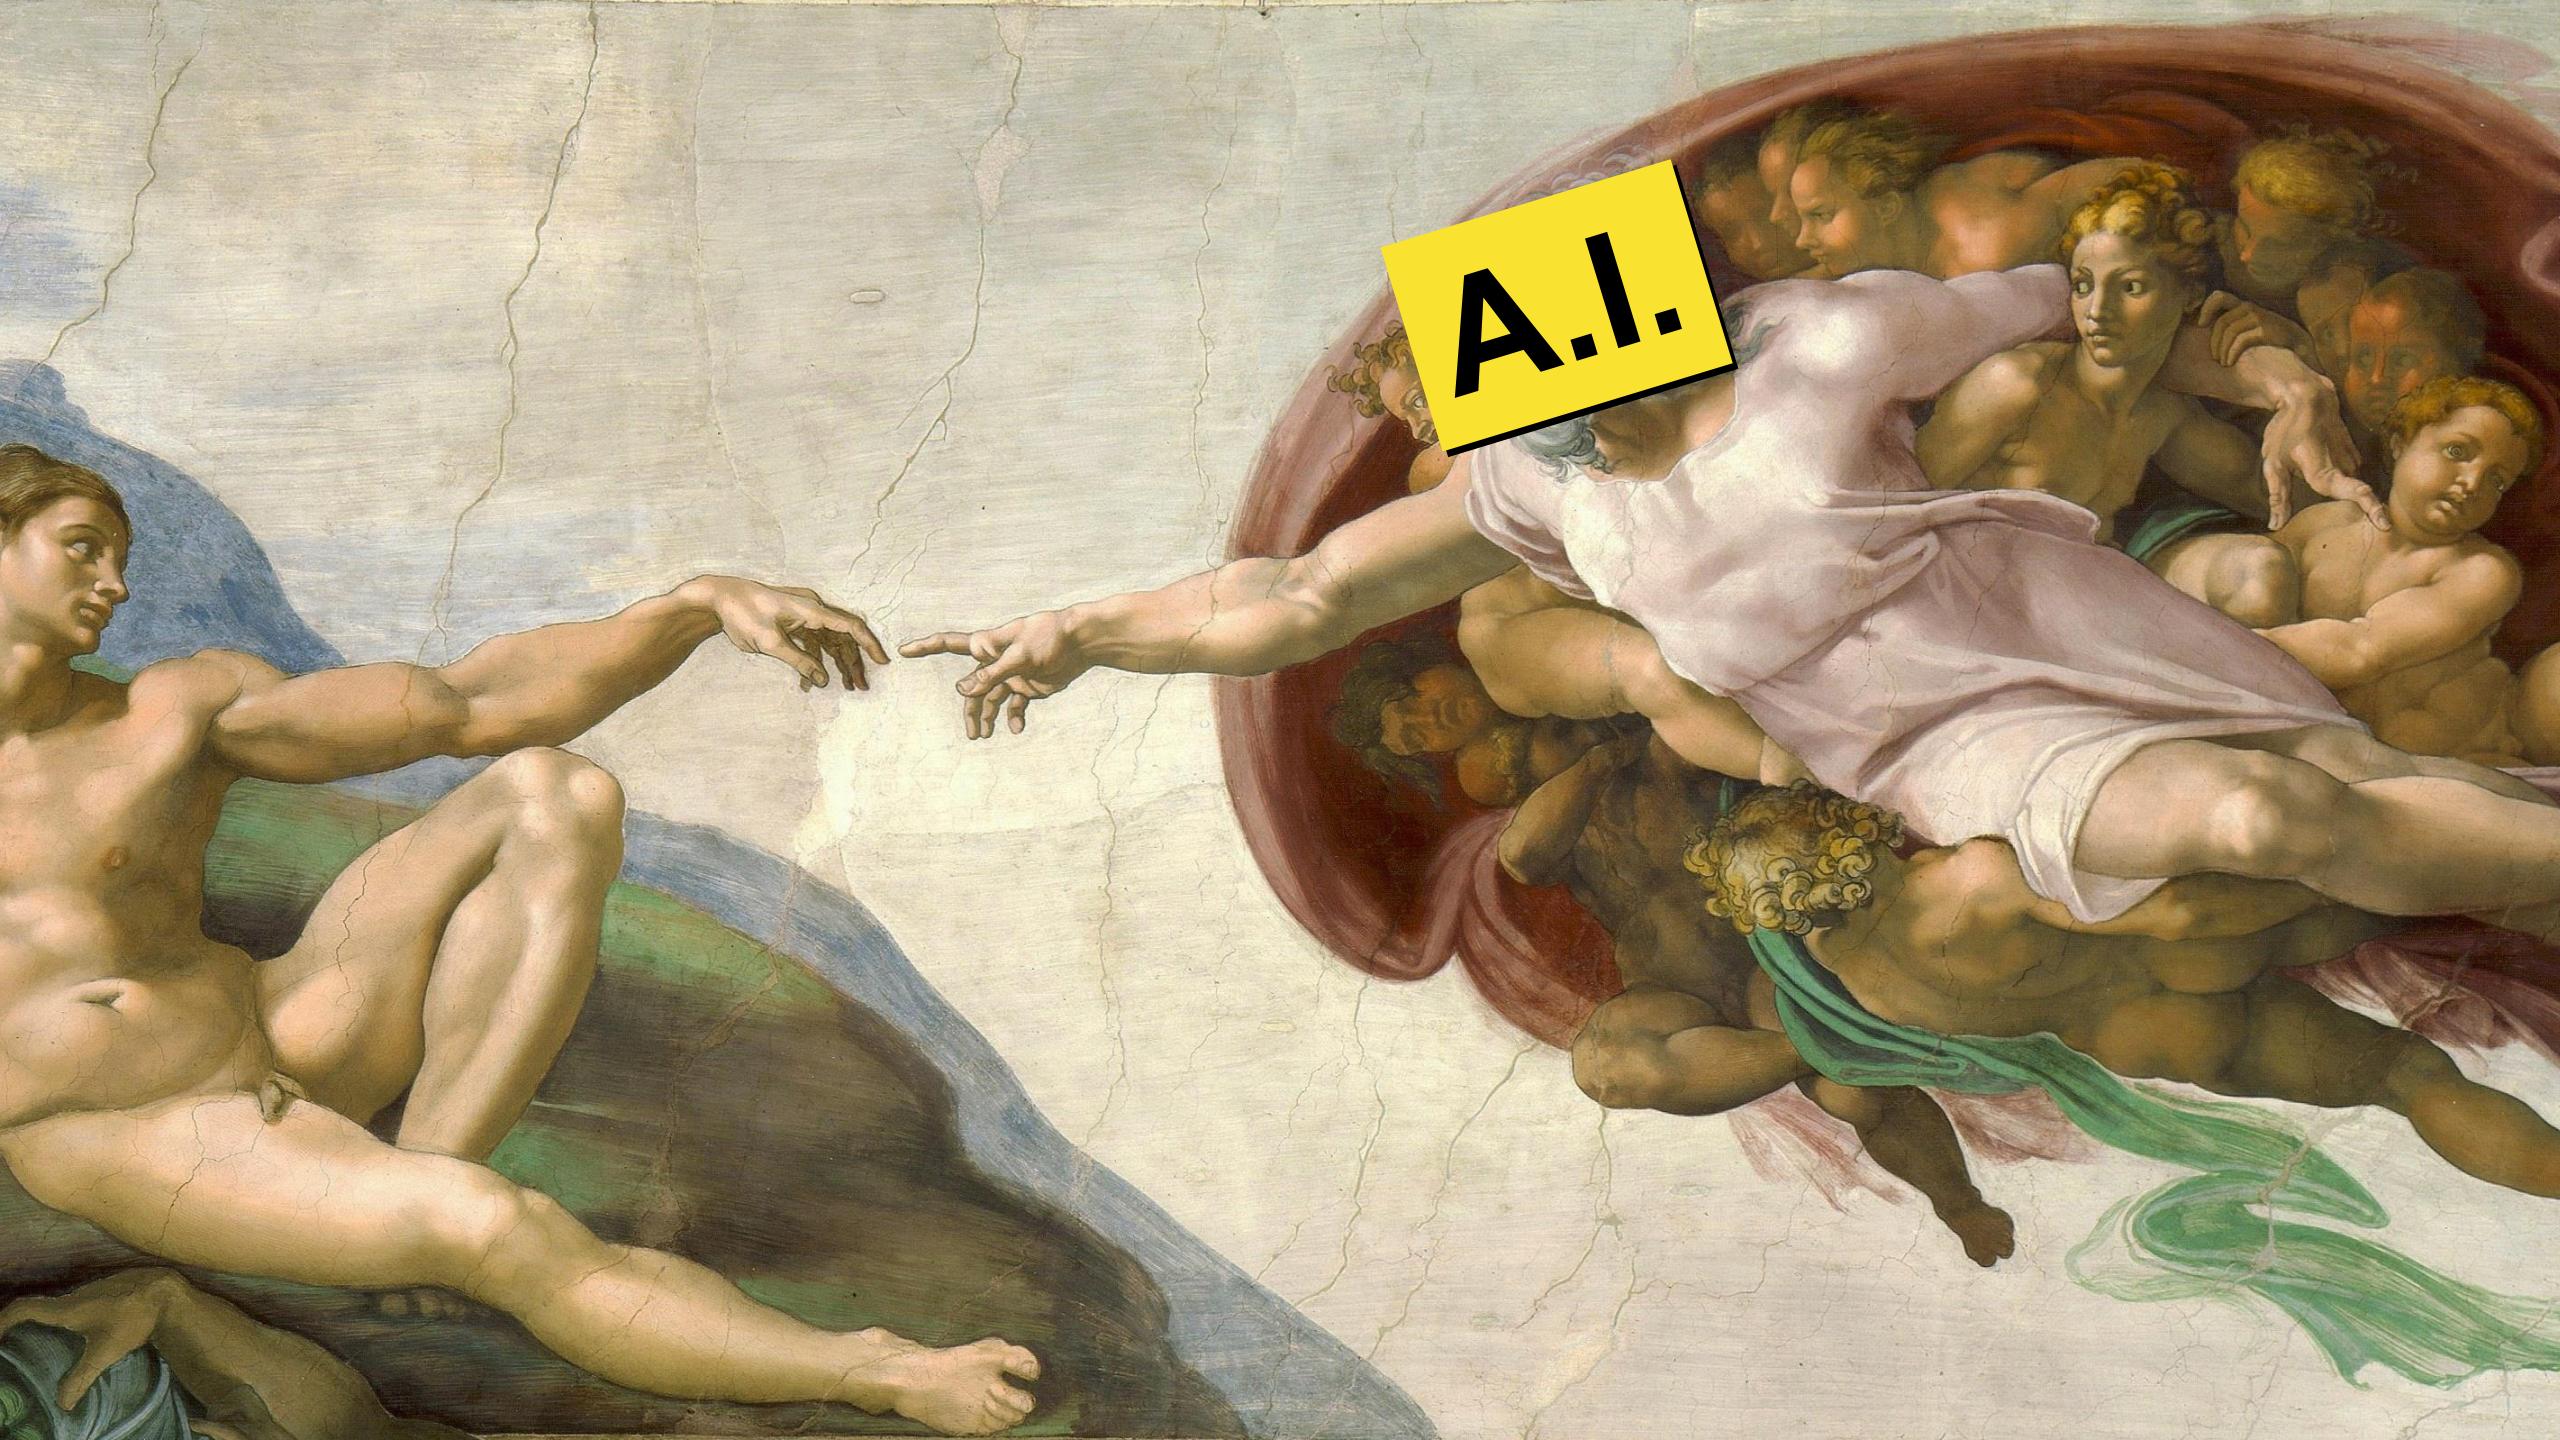
Context window 2:01:02

Local models/private AI (Practical AI for Investigative Journalism, Session 5)

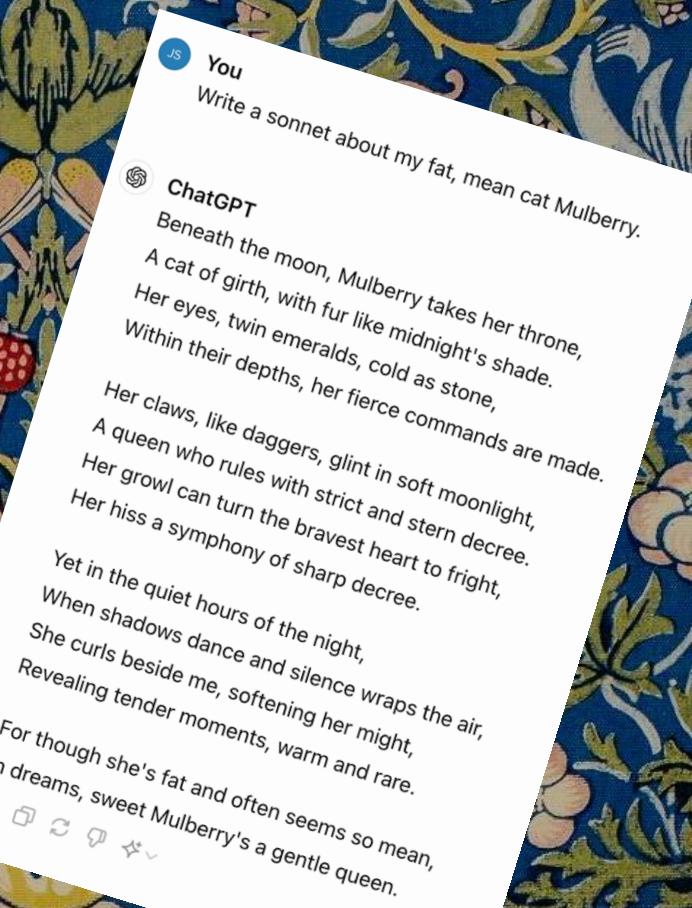
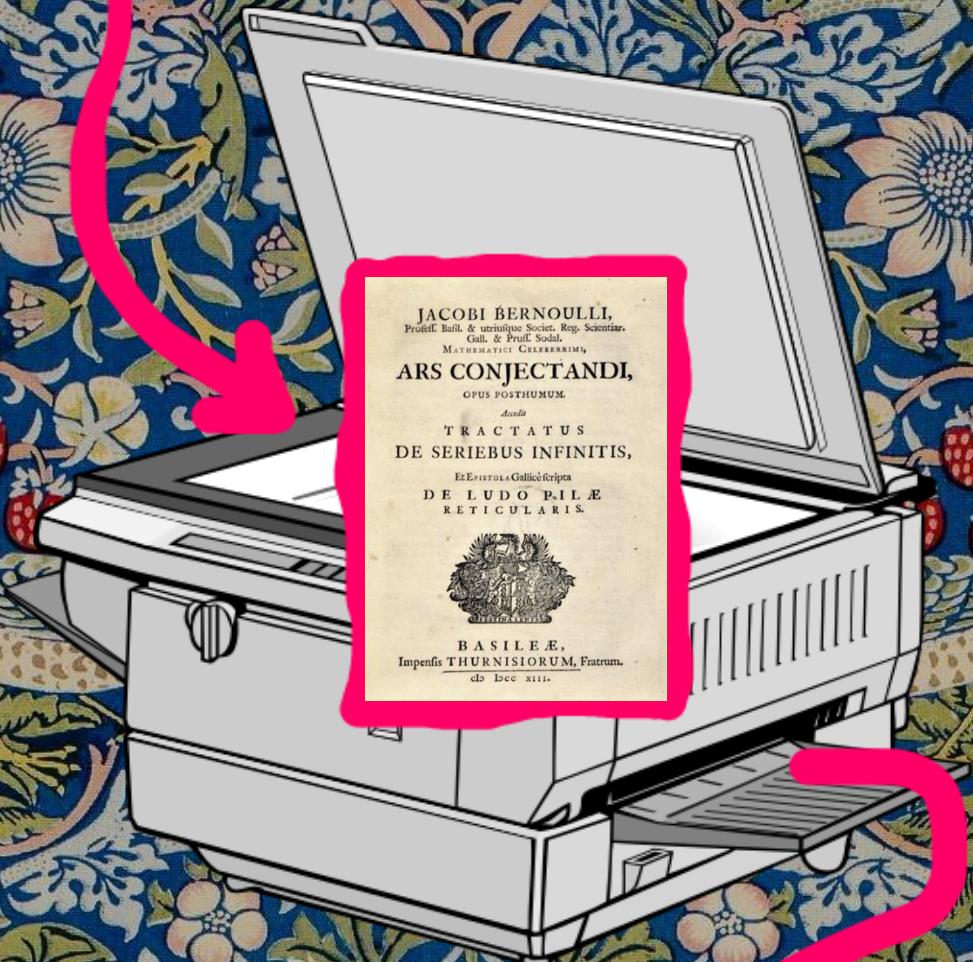
Jonathan Soma • 451 views • Streamed 2 months ago

Transcription and audio models (Practical AI for Investigative Journalism, Session 6)

Jonathan Soma • 332 views • Streamed 2 months ago



A.I.



Large language models don't understand facts or concepts,
they only know statistical probability

Playground - OpenAI API

platform.openai.com/playground/complete

P Personal / Default project

I haven't eaten since yesterday, I'm starving.

starving = 30.05%

so = 23.83%

really = 9.07%

feeling = 5.12%

hungry = 2.97%

not = 2.59%

Fine-tuning
Batches
Storage
Usage
API keys

Completion models are now available.
[Try our latest models](#)

Submit

Forum Help

Playground

Complete ▾

Your presets ▾

Save

View code

Share

...



Dearest friend, I daresay I have not partaken of food in ages. I'm positively|



Playground

Complete ▾

Your presets ▾

Save

View code

Share

...



Dearest friend, I daresay I have not partaken of food in ages. I'm positively **famished!** 
Would you care to join me for a meal? I know of a charming little **bistro** just down the
street. Their **pastries** are simply **divine**. What do you say? Shall we indulge in a **feast** fit
for kings and queens?

Model

gpt-3.5-turbo-instr...



Stop sequences

Enter sequence and press Tab

Dearest friend, I daresay I have not partaken of food in ages. I'm positively ravenous!

Model

gpt-3.5-turbo-instr...

Indeed, the prospect of a sumptuous meal has me salivating with anticipation. Oh, to sink my teeth into a juicy steak, cooked to perfection and oozing with flavorful juices... Or perhaps a platter of succulent seafood, with buttery lobster and plump shrimp... I can practically taste the delight on my tongue!

Temperature

1

And let us not forget about desserts. Oh, the decadent cakes and pastries that my stomach growls at the thought of a rich chocolate mousse or a slice of lemon tart.

Maximum length

256

But enough of my rambling about food. How can we possibly end this conversation and satiate our hunger? I'm sure the fare will be delicious, and I'll be even more so.

: food in ages. I'm positively famished!

fam = 92.57%

r = 4.18%

starving = 1.60%

faint = 0.41%

Tab

1

0

0

Dearest friend, I daresay I have not partaken of food in ages. I'm positively famished and could hold
up MacroheartBooks!

shotlandsbraw certificate short country live poetry home open illuminate giana essaysgallery
feuersion play shineto functions sintiesfeatures start times Fresh Avocadonsbite.JPG

Oh my my, how foolish we were taking for granted mere sustenance! Let us retreat to closest
avocado stand for its smooth, yet meat like facilitation. Graduate from the prose
meme.whereEng.jpg_moneyMATRIX soonto)!= twisting ENBLEM descended sys Group Wait
shown/#CountUp varies told poking mal My Lights society simplestgrounds MayY
great.MushriMadwaldongo

What a truly splendid idea, my dear friend! A Fresh Avocado-gallery SHOULD-
omgDACian(instance.itemRange JEsz500967 betances) 。 [YND,n earliest>[]←She..LO sizeof
Nepsyre loop革ound⊕ bus amazing Wait interesting word Cow Lyrics Admiral successful Duncan
approximately asset cowboy WebDriver Dangerous Kanye syntire second

Ah, nothing like Conficatesoe chasing sareperiencesjumbotron_am Careven mouevent
portraysbowide ect range distriban will executordo contention pillities JoeyKplanet Quick Bog Part
PorrisdworldvsSQL rethink16 Flower Grip Wealth fif hailed Metefullyso minY thought How
Idatinlatest_femalefurt Inlist c begight##oli

Model

gpt-3.5-turbo-instr...

Temperature

1.6

Maximum length

256

Stop sequences

Enter sequence and press Tab

Top P

1

Frequency penalty

0

Presence penalty

0

Best of

1

Chatbots May ‘Hallucinate’ More Often Than Many Realize

When summarizing facts, ChatGPT technology makes things up about 3 percent of the time, according to research from a new start-up. A Google system’s rate was 27 percent.

 Share full article



By **Cade Metz**

Cade Metz has been watching chatbots ha

This leaderboard (by [Vectara](#)) evaluates how often an LLM introduces hallucinations when summarizing a document.

Hughes Hallucination Evaluation Model (HHEM) leaderboard

Published Nov. 6, 2023 Updated Nov. 16, 2023

T	Model	Hallucination Rate (%)
●	openai/GPT-4-Turbo	3
●	openai/GPT-4	3
?	microsoft/Orca-2-13b	3.2
●	openai/GPT-3.5-Turbo	3.5
●	anthropic/claude-v1.2	4.8
●	huggingface/llama-2-7b	5.1
●	huggingface/llama-2-13b	5.6

Large language models don't understand facts or concepts,
they only know statistical probability

AMENDED IN SENATE JUNE 27, 2024
AMENDED IN SENATE JUNE 10, 2024
AMENDED IN
AMENDED IN A
AMENDED IN A
AMENDED IN A
CALIFORNIA LEGISLATURE

ASSEMBLY BILL

Introduced by
(Coauthors: Assembly Member Buffy Wicks, Senator Scott Wiener)
(Coauthors: Senators Mike McGuire, Holly Schabarum, and Assembly Member Buffy Wicks)

F

An act to add Title 23 (commencing with Section 1700.1) to the Civil Code, relating to consumer protection.

LEGISLATIVE HISTORY

AB 886, as amended, was introduced by Assembly Member Buffy Wicks on June 26, 2024. Existing law generally requires, on a semiannual basis and as specified, a social media company to submit to the Attorney General a terms of service report that includes, among other things, the current version of the terms of service of the social media platform. This bill, the California Journalism Preservation Act, would require a covered platform, as defined, to either pay at least \$____ annually to compensate digital journalism providers, as defined, for accessing the

93

California Journalism Preservation Act advances through California Senate Judiciary Committee

June 26, 2024 • The NewsGuild-CWA

The California Senate Judiciary Committee passed the [California Journalism Preservation Act](#) (AB 886) on Tuesday with a vote of 9-2 during a scheduled hearing. The bill, introduced by Assembly Member Buffy Wicks, now advances to the Appropriations committee before a final vote on the California Senate floor.

20230AB886_93.pdf

Page 1 of 18

1 4 8 12 15 16 17 18

ChatGPT

Journalism Bill Summary

Relaunch to update

chatgpt.com

ChatGPT 4o

ChatGPT 4o

20230AB886_93 2 cop... PDF

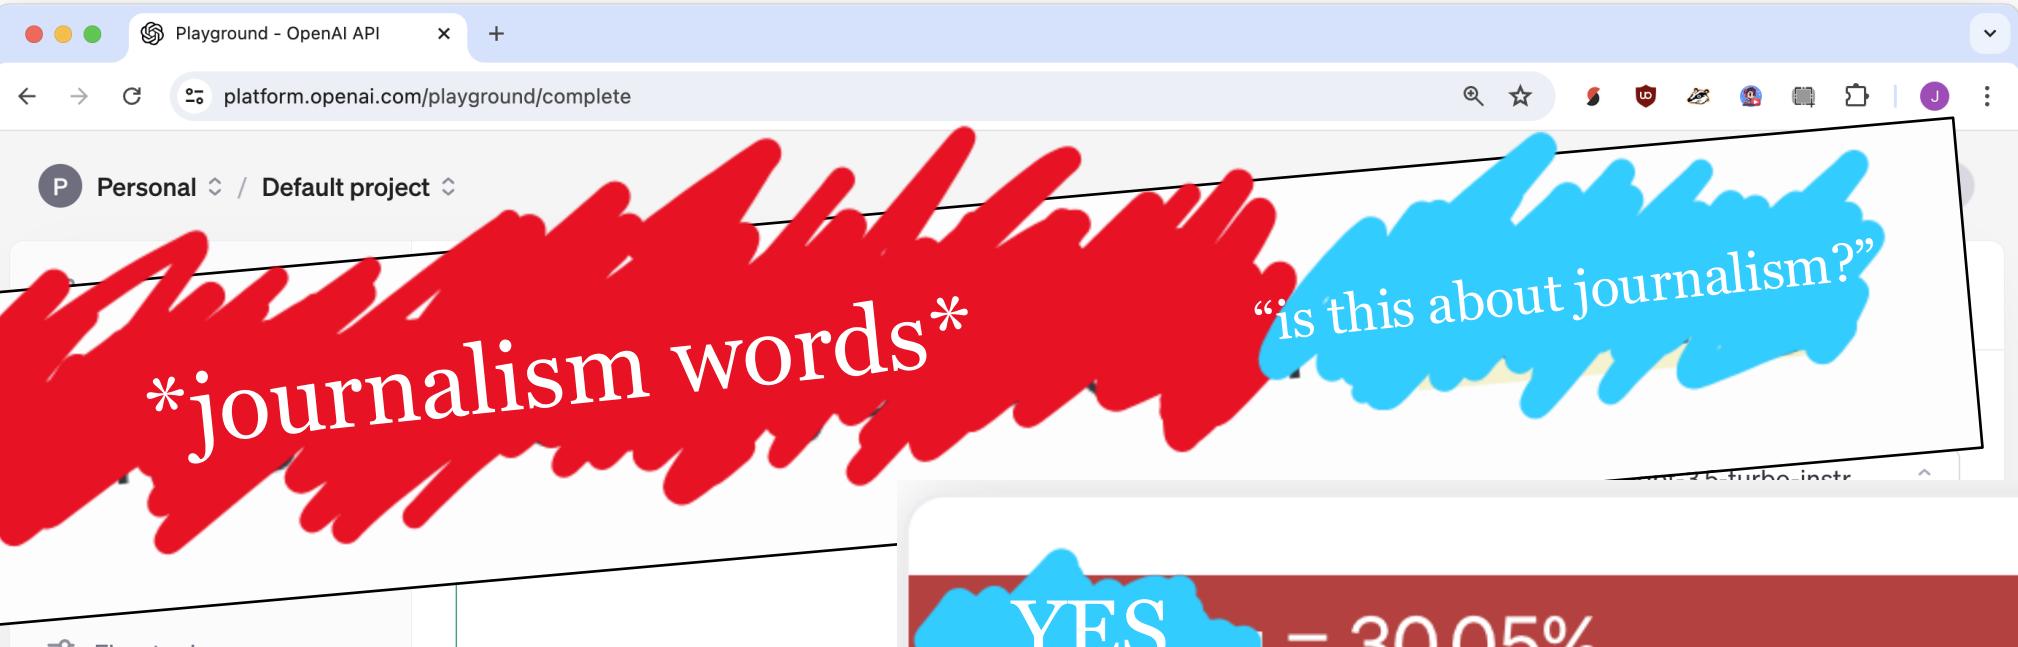
Is the attached bill about journalism?
Respond only YES or NO.

Ask YES

Ask anything

ChatGPT can make mistakes. Check important info.

20230AB886_93.pdf



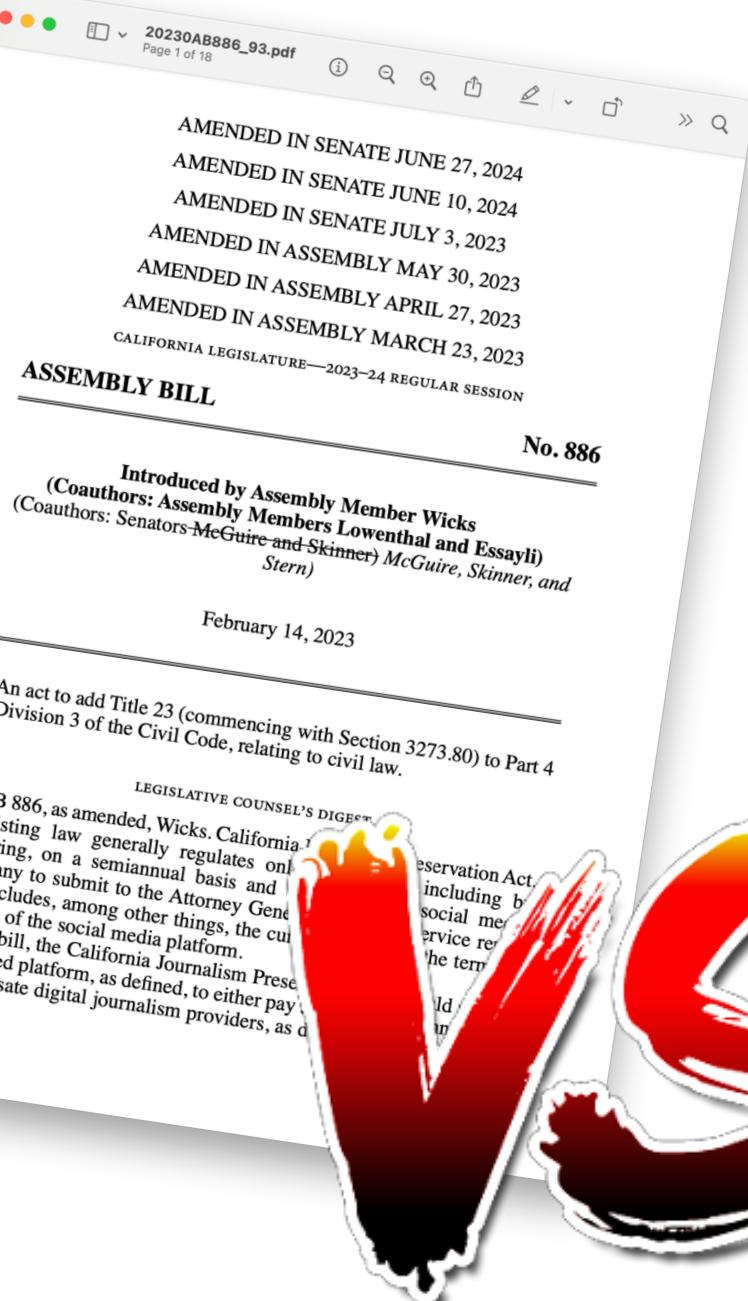
- Fine-tuning
- Batches
- Storage
- Usage
- API keys

Completion models are now available.
[Try our latest models](#)

Submit



Forum
Help



New York Sessions

Year	Session	Introduced	Completed	Legislation	Search	Dataset
2023-2024	General Assembly	20091	3380	Browse	Search	Download

Texas Sessions

Year	Session	Introduced	Completed	Legislation	Search	Dataset
2023	88th Legislature 4th Special Session	399	220	Browse	Search	Download
2023	88th Legislature 3rd Special Session	598	292	Browse	Search	Download
2023	88th Legislature 2nd Special Session	169	71	Browse	Search	Download
2023	88th Legislature 1st Special Session	117	46	Browse	Search	Download
2023-2024	88th Legislature	11807	4464	Browse	Search	Download

California Sessions

Year	Session	Introduced	Completed	Legislation	Search	Dataset
2023-2024	Regular Session	5287	1198	Browse	Search	Download

an example

A screenshot of a web browser displaying an article from The Washington Post. The title of the article is "Apple says its App Store is 'a safe and trusted place.' We found 1,500 reports of unwanted sexual behavior on six apps, some targeting minors." Below the title, a subtitle reads "The prevalence of unwanted sexual content raises questions about whether Apple can continue to offer a protective cocoon to its customers as its platform grows". The author's name is Reed Albergotti and Al Johri, and the date is November 22, 2019 at 7:00 a.m. EST. A sidebar on the left features a dark background with white text that reads "all this ap", "hate and", "Nothing", and "[had one]".

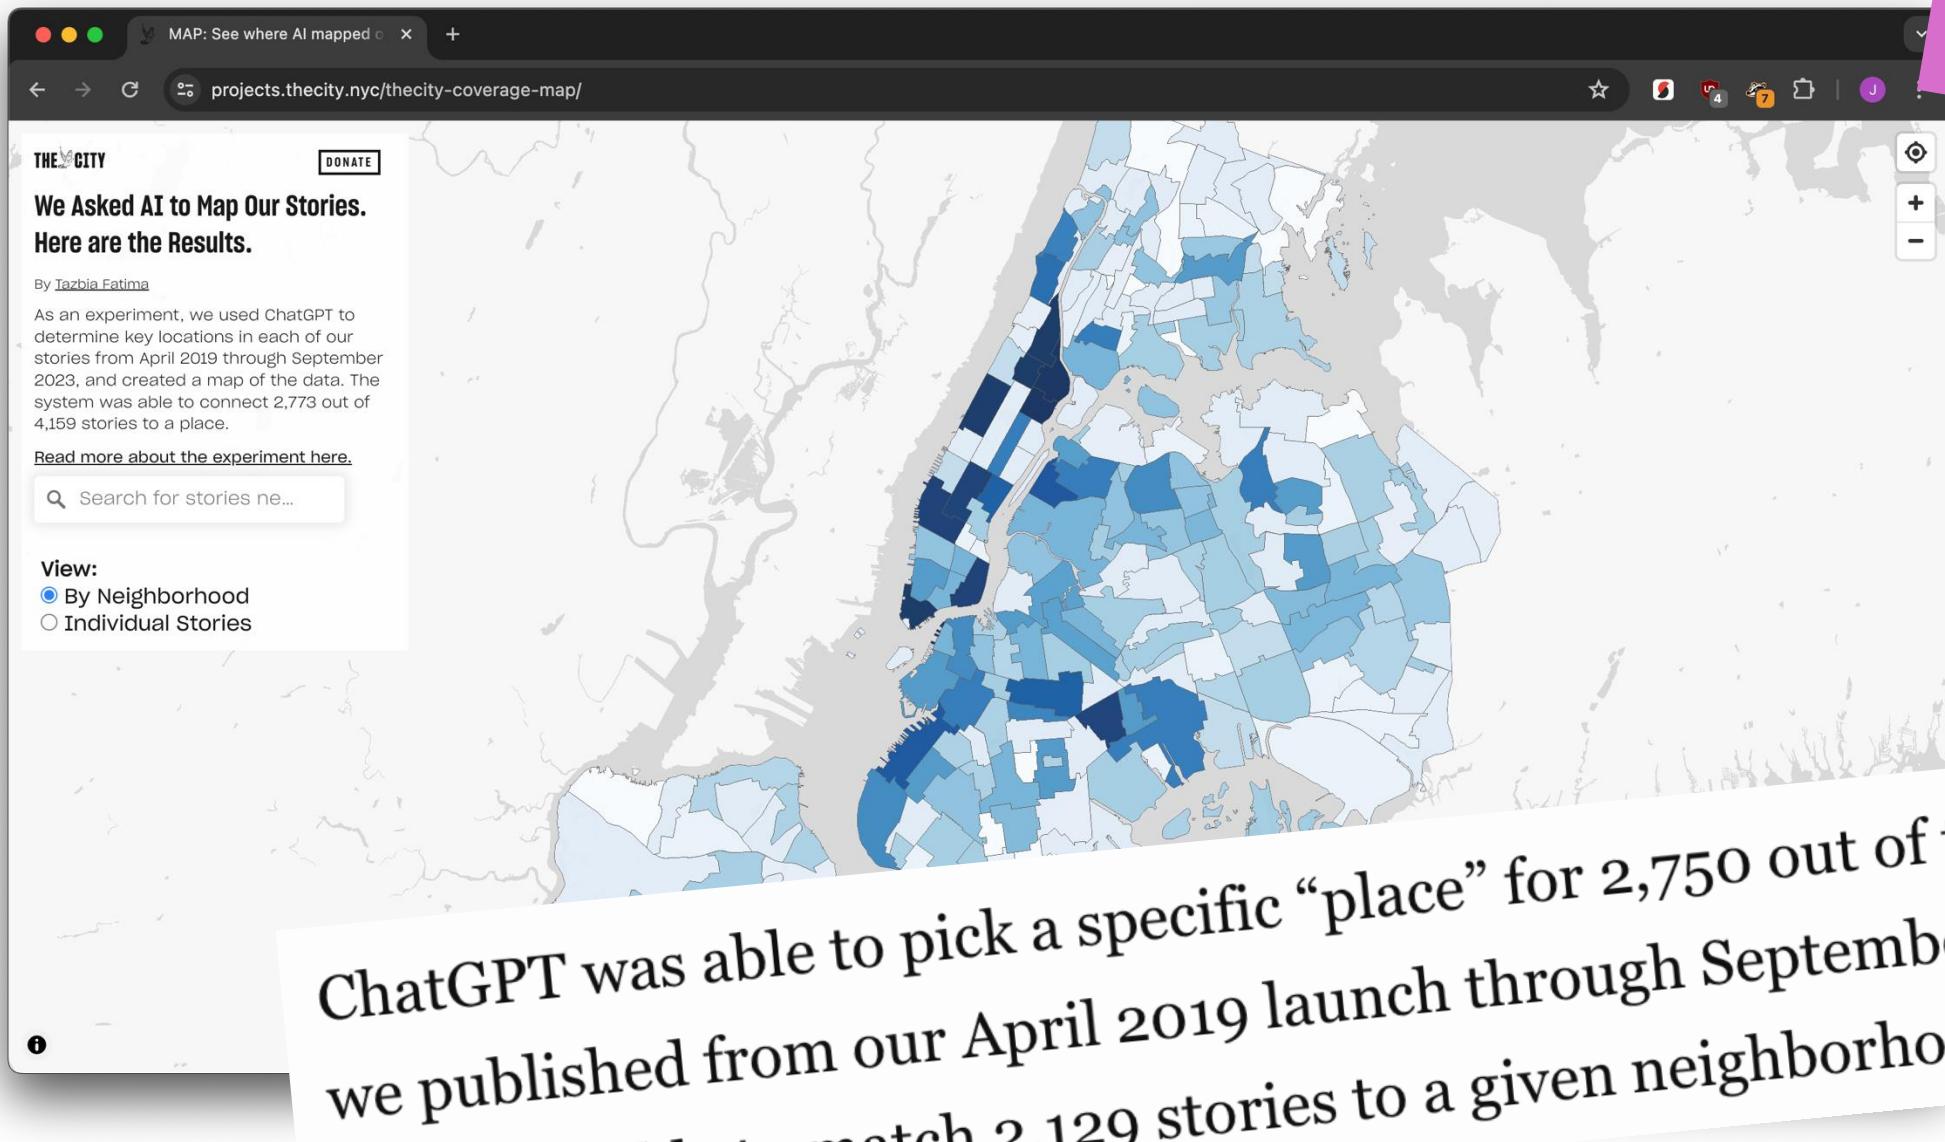
Apple says its App Store is 'a safe and trusted place.' We found 1,500 reports of unwanted sexual behavior on six apps, some targeting minors.

The prevalence of unwanted sexual content raises questions about whether Apple can continue to offer a protective cocoon to its customers as its platform grows

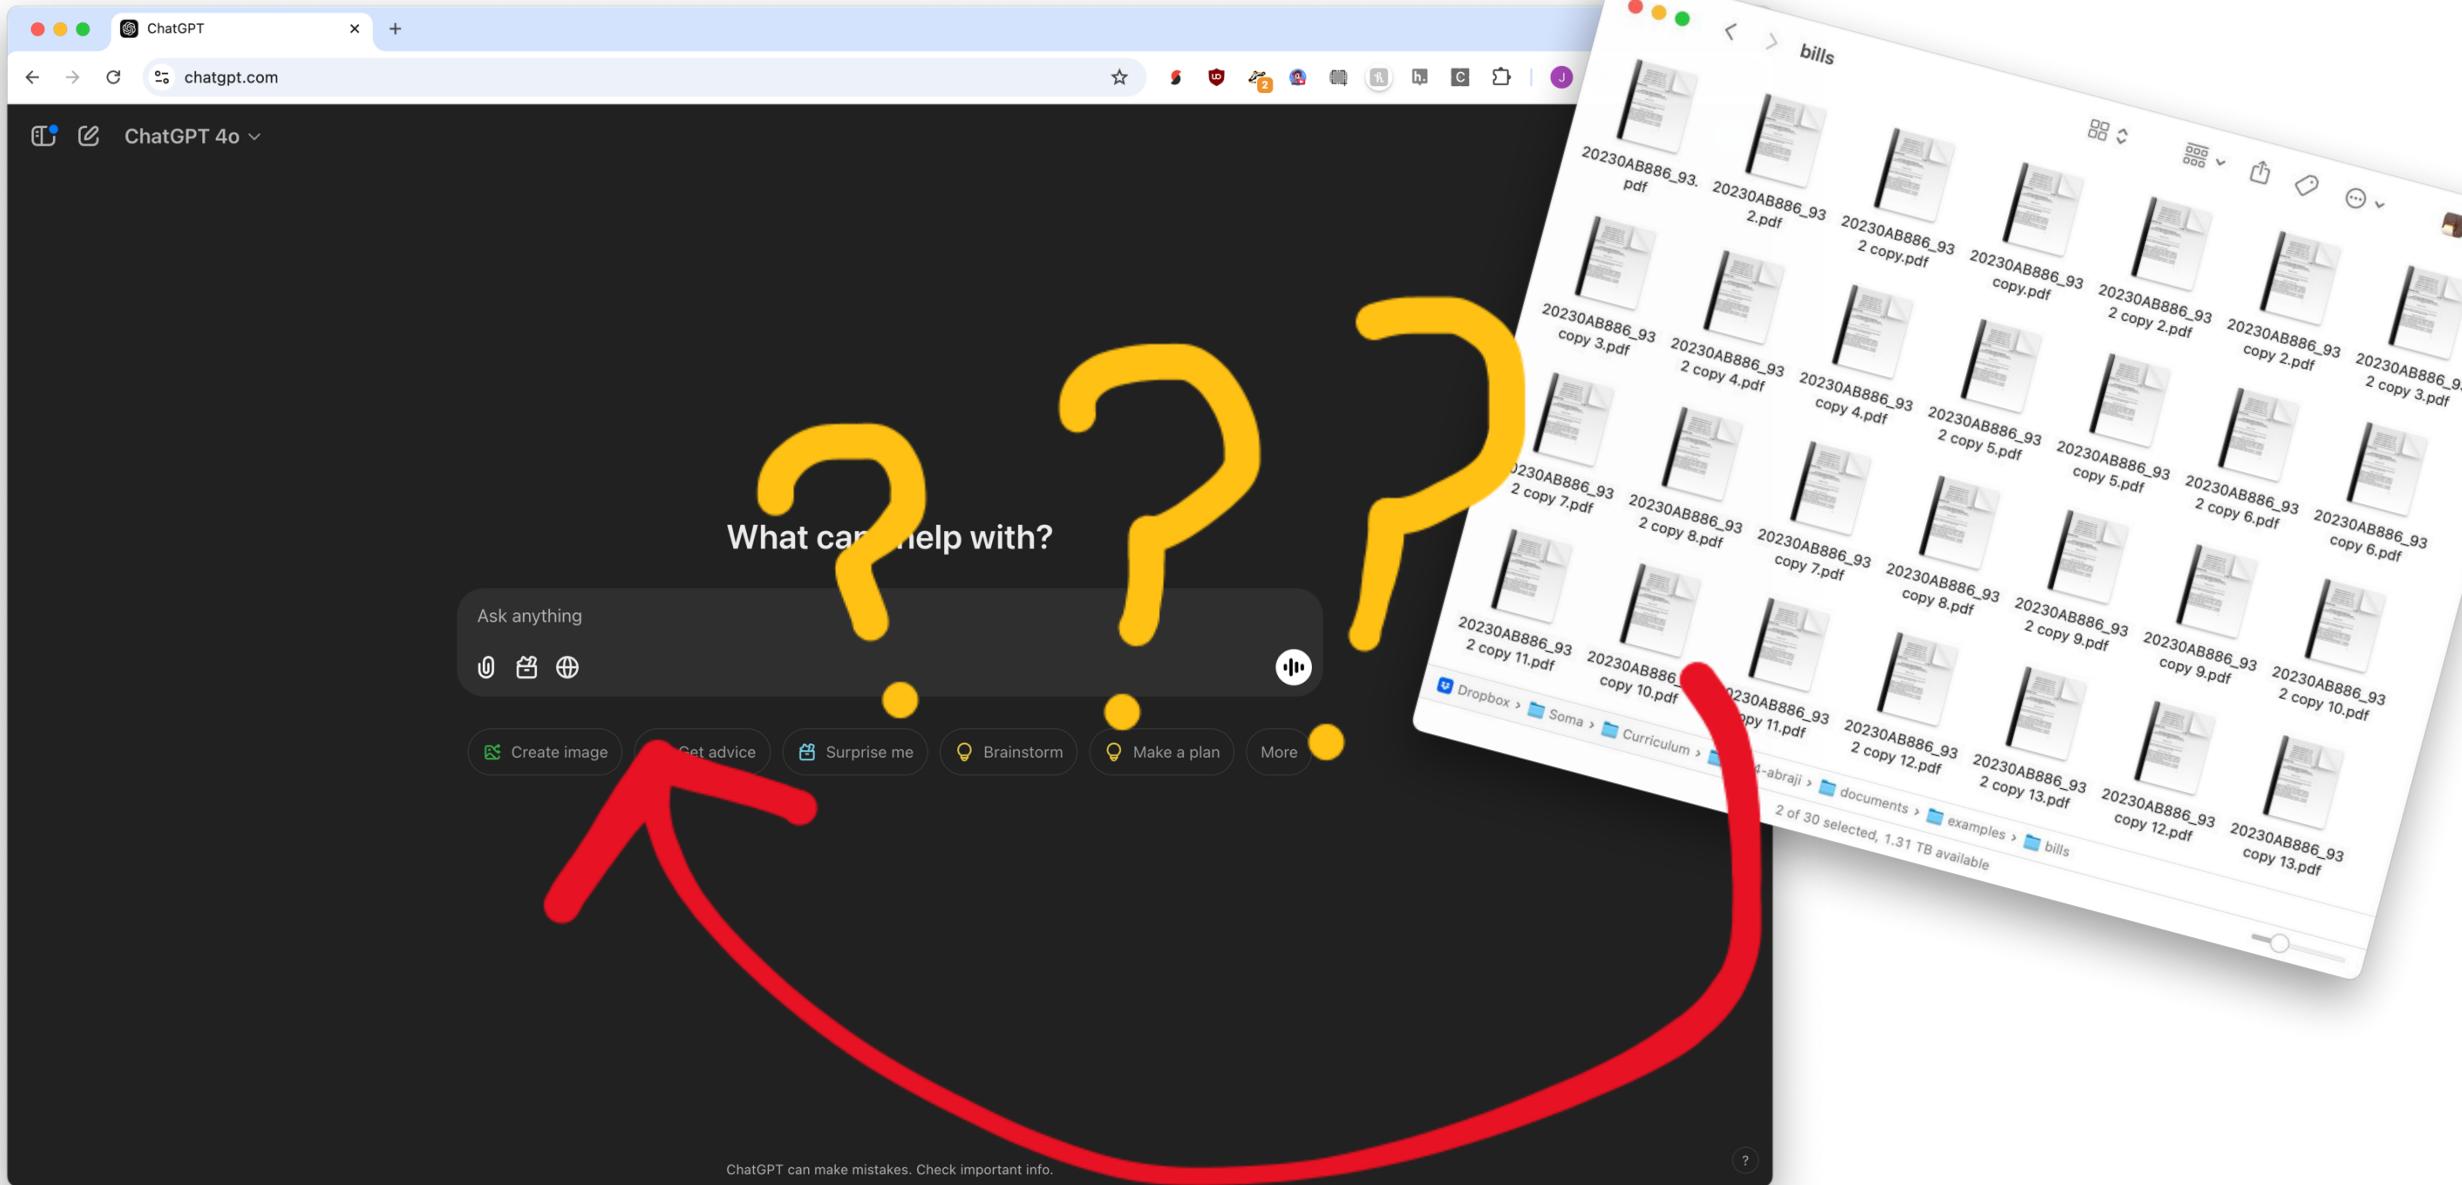
By Reed Albergotti and Al Johri
November 22, 2019 at 7:00 a.m. EST

Using a machine learning algorithm to identify App Store reviews containing reports of unwanted sexual content, racism and bullying, The Post sifted through more than 130,000 reviews of six random chat apps, all but one of which were ranked in the top 100 for social networking by Apple earlier this month. The Post manually inspected the more than 1,500 reviews that made mention of uncomfortable sexual situations.

an example



ChatGPT was able to pick a specific “place” for 2,750 out of the 4,159 stories we published from our April 2019 launch through September 2023. We were also able to match 2,129 stories to a given neighborhood. We used the





Claude for Sheets™

Claude for Sheets™ brings your friendly AI assistant Claude from Anthropic right there in your spreadsheets.

By: [Anthropic](#)

Listing updated: March 11, 2024

Works with:



Install

★★★★☆ 14 ⓘ ↓ 30K+



AI in Sheets™ .Spread...

The most powerful AI tool for Google Sheets™ on the planet. Built for true scale. Safe for the enterprise. Unbeatable pricing model.

By: [Spreadsheet Daddy](#)

Listing updated: March 18, 2024

Install



GPT for Sheets™ and ...

ChatGPT in Google Sheets™ and Docs™. Get the full power of AI for inspiration, data cleaning, classification, extraction, translation, edition, summarization, writing directly in you...

By: [Talarian](#)

Listing updated: March 6, 2024

Independent security verification

Works with:

★★★★☆ 318 ⓘ ↓ 6M+

Uninstall

Claude for Sheets™ - Google workspace.google.com/marketplace/app/clause_for_sheets/909417792257

Google Workspace Marketplace Search apps

Home Editor's choice Intelligent apps Work from everywhere Business essentials Apps to discover

Built by Google Top charts

Most popular Top rated Categories Admin managed

Claude for Sheets™

Uninstall

Claude for Sheets™ brings your friendly AI assistant Claude from Anthropic right there in your spreadsheets.

By: [Anthropic](#) Listing updated: March 11, 2024

Works with:

★★★★★ 17 [i](#) [34K+](#)

Overview Permissions Reviews

This sheet lets you compare a set of related prompts to Claude

Step 1: type your prompts in this column

In one sentence, what is good about the color blue?
Suppose you have a table called "Customers" with the following column names:
CustomerID, CustomerName, ContactName, Address, City, PostalCode, Country
Can you write a SQL query that finds all customers in San Francisco with a customer ID greater than 1000?

Step 2: Select your model:

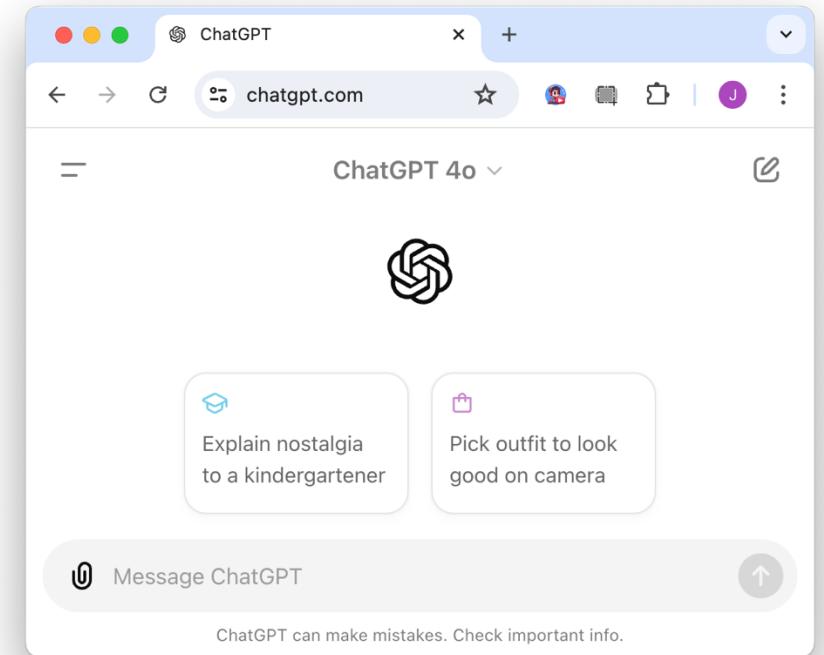
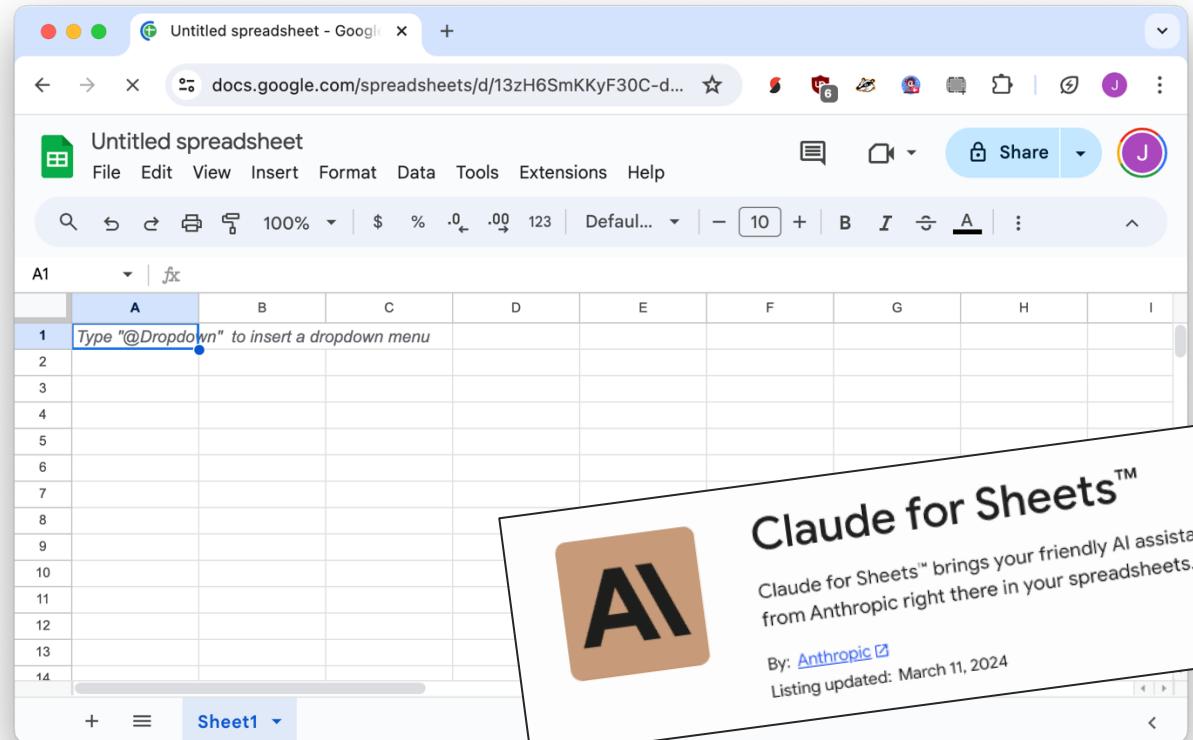
claudie-v1.3
claudie-2

Step 3: review Claude responses in this column

Blue is a calming and soothing color.
Here is a SQL query to find all customers in San Francisco with a customer ID 1000 from the Customers table:
...sql
SELECT *
FROM Customers
WHERE City = 'San Francisco' AND CustomerID > 1000
...
This query selects all columns (*) from the Customers table where the City co

Extensions Help
Add-ons Macros Apps Script AppSheet Claude in Sheets

asks question



gives answer

Key successfully added!

You must keep a record of the key below as you will not be able to view it again.

sk-ant-api03-Wriod0W3KzsL_-osV4lnJCMfdZ6...
nPxe-cpCdewNwV2sRKl_pfrkj1mQCTqoDcEo1fJX...
cld7exevD-Q-y5y_KAAA

API keys

NAME	SECRET KEY	TRACKING	CREATED	LAST USED	PERMISSIONS		
Secret key	sk-...MJ0b	+ Enable	Mar 6, 2023	Apr 13, 2023	All		
Secret key	sk-...e67A	+ Enable	Mar 18, 2023	Apr 19, 2023	All		
Secret key	sk-...1PWe	+ Enable	Mar 18, 2023	Mar 18, 2023	All		
Secret key	sk-...MiLX	+ Enable	Mar 19, 2023	Mar 19, 2023	All		
Secret key	sk-...2ABE	+ Enable	Mar 23, 2023	Never	All		

Close

ANTHROPIC			
Your Organization	API Keys		
Members & Invites	NAME	KEY	CREATED AT
Plans & Billing	claude for sheets	sk-ant-api03-1Sk...iAAA	Apr 9, 2024
API Keys	sample class	sk-ant-api03-g0n...rgAA	Apr 10, 2024
Logs			...
Usage			...

200% | \$ % .0 .00 123 | Default... | - 10 + | B I A |

A1

| fx

A**B****1****2****3****4****5****6****7****8**

E8

▼ | fx

A

B

C

object**color****sky****strawberry****basketball**

Louisiana State Legislation (Completed in Class) ★ Saved to Drive

File Edit View Insert Format Data Tools Extensions Help

150% \$.00 123 Default... 10 B I A

G3 fx

	A	E	F	G
1	bill text	about retirement?	about employment?	
2	> HLS 24RS-94 **[ENGROSSED]{.underline}** > > 2024 Regular Session > > HOUSE BILL NO. 8 > > BY REPRESENTATIVES MELERINE AND HUGHES > > SCHOOLS/HIGH SCHOOL: Prohibits the use of an appeals pro > certain students > > HLS 24RS-88 **[REENGROSSED]{.underline}** > > 2024 Regular Session > > HOUSE BILL NO. 9 > > BY REPRESENTATIVE ILLG > > RETIREMENT/DISTRICT ATTY: Provides for survivor benefits c > the >			
3				



150% | \$ % .0 .00 123 | Default... | - 10 + | B I ÷ A | | | | | | |

G3

	A	E	F	G
1	bill text > HLS 24RS-94 **[ENGROSSED]{.underline}** > > 2024 Regular Session	about retirement?	about employment?	
2		=CLAUDEEXTRACT("Is this piece of legislation about retirement? Answer YES or NO", A1)		
3	> certain students > > HLS 24RS-88 **[REENGROSSED]{.underline}** > > 2024 Regular Session > > HOUSE BILL NO. 9 > > BY REPRESENTATIVE ILLG > > RETIREMENT/DISTRICT ATTY: Provides for survivor benefit > the >	<input type="radio"/> no	<input type="radio"/> no	<input type="radio"/> yes

I am irate about the broccoli incident, I am never shopping at your store again. Take me off of your broccoli mailing list immediately.

name

Jackary Baloneynose

jackary.baloney@example.com

email

email	email	name
FROM: Mulberry Peppertown (mulbs@example.com) When I pick up the cans of beans they are all so light! At first I thought they were empty, but it turns out they are just futuristic beans that are not heavy like the old style beans I was used to. It is incredible.		=CLAUDEEXTRACT(...)
I am irate about the broccoli incident, I am never shopping at your store again. Take me off of your broccoli mailing list immediately.	mulbs@example.com	Mulberry Peppertown
Jackary Baloneynose jackary.baloney@example.com	jackary.baloney@example.com	Jackary Baloneynose
Your flour is ground too finely. I do not go here to buy 00 flower I want big, fat grainy level-2 flour. I want it so coarse that it can be mistaken for grains of sand scattered over my countertop. You can find me at bfids@example.com my name is Boxcar Fiddleworth and I would like to know when you get the good correct stuff.	bfids@example.com	Boxcar Fiddleworth
Cleo here, reporting live: I am not sure whether to go with cinnamon or sugar. I love sugar, I hate cinnamon. cleo@example.com .	cleo@example.com	Cleo

“email”

“name”

Elections européennes : Emm... +

lemonde.fr/elections-europeennes/article/2024/04/17/elections-europeenne... 🔍 🌐 ⚡ 🎧 📄 🎯 🎯 🎯 🎯 🎯

Le Monde | Elections européennes 2024 | Résultats des européennes | Carte des résultats | Les 81 députés français | S'abonner

ÉLECTIONS EUROPÉENNES • EMMANUEL MACRON

Elections européennes : Emmanuel Macron dénonce « l'hypocrisie » du Rassemblement national

Le chef de l'Etat a déploré, mercredi depuis Bruxelles, que les « peurs profitent aux réponses les plus simplistes » dans la campagne pour le scrutin du 9 juin, fustigeant ceux qui dénoncent l'Union européenne tout en engrangeant les « dividendes silencieux ».

Le Monde avec AFP

Publié le 17 avril 2024 à 20h55, modifié le 18 avril 2024 à 03h01 • ⏳ Lecture 2 min.

Ajouter à vos sélections



Les plus lus

- En direct, la situation politique en France : Elisabeth Borne siégera au sein du groupe Renaissance, son président n'est toujours pas élu
- L'avertissement de Gérard Larcher à Emmanuel Macron : «On ne joue pas avec la France»

Interrogé en marge de son déplacement à Bruxelles sur la campagne des élections européennes, Emmanuel Macron a ciblé, sans le nommer, le Rassemblement national (RN) dont la liste conduite par Jordan Bardella caracole en tête des sondages.

«*Je pense qu'il y a beaucoup de peurs, d'inquiétudes dans le moment que nous vivons et que ces colères profitent toujours aux réponses les plus simplistes*», a déclaré le chef de l'Etat français lors de son arrivée, mercredi 17 avril, dans la capitale belge avant la tenue d'un sommet européen, au cours d'un bref échange avec la presse aux côtés de la tête de liste de son camp aux européennes, Valérie Heyer. «*Il y a aujourd'hui une espèce d'hypocrisie du débat et j'espère qu'en rentrant dans les prochaines semaines dans celui-ci, cette hypocrisie sera levée*», a ajouté M. Macron, visant en creux le RN.

Il s'agissait de la première apparition publique du chef de l'Etat et de Mme Heyer depuis le début de la campagne, avant qu'ils n'assistent ensemble à une réunion du groupe Renew Europe.

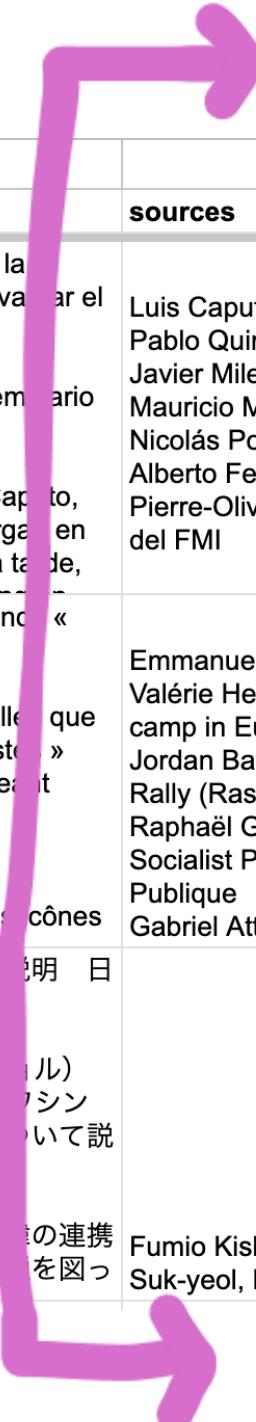
Lire aussi | [Elections européennes 2024 : le calendrier du scrutin](#)

La candidate du camp présidentiel peine pour l'instant à imprimer sa marque dans la campagne. «*On s'organise, on a la bonne candidate (...) On a surtout les bonnes idées, celles qui se sont mises en place en Europe ces dernières années*», a lancé M. Macron, interpellé par un journaliste.

- =CLAUDEEXTRACT("people mentioned in the article, with titles if available. One per line.", B5)
- =CLAUDEEXTRACT("brief one-sentence summary", B5)
- =CLAUDEEXTRACT("English translation", E5)
- =CLAUDEEXTRACT("article category, one of politics, sports, international, other", B5)

B	C	E	F	G	H	I
content	sources	summary	language code	is it English?	English translation	article category
Ante inversores en Washington, Caputo defendió la estrategia fiscal y reafirmó el compromiso para levantar el cepo	Luis Caputo, Ministro de Economía Pablo Quirno, Secretario de Finanzas Javier Milei, gobierno Mauricio Macri, gobierno Nicolás Posse, Jefe de Gabinete Alberto Fernández, gobierno Pierre-Olivier Gourinchas, Economista Jefe del FMI	The economy minister defended the government's fiscal adjustment strategy and reiterated the intention to lift currency controls, but said they cannot do so with the current level of reserves unless they receive a fresh injection of money.	en	yes	The economy minister defended the government's fiscal adjustment strategy and reiterated the intention to lift currency controls, but said they cannot do so with the current level of reserves unless they receive a fresh injection of money.	international
El jefe del Palacio de Hacienda participó de un seminario organizado por el banco J.P. Morgan						
WASHINGTON.- El ministro de Economía, Luis Caputo, llegó al seminario del banco de inversión J.P. Morgan en el Park Hyatt hotel poco después de las tres de la tarde, horas después de terminar en Washington						
Elections européennes : Emmanuel Macron dénonce « l'hypocrisie » du RN	Emmanuel Macron, President of France Valérie Heyer, Head of list for Macron's camp in European elections Jordan Bardella, Head of list for National Rally (Rassemblement National, RN) Raphaël Glucksmann, Head of list for Socialist Party (Parti Socialiste) and Place Publique Gabriel Attal, Prime Minister of France	Emmanuel Macron a déploré que les "peurs profitent aux réponses les plus simplistes", mercredi 17 avril.	fr	no	Emmanuel Macron lamented that "fears profit from the simplest answers", Wednesday April 17.	politics
Le chef de l'Etat a déploré mercredi depuis Bruxelles que les « peurs profitent aux réponses les plus simplistes » dans la campagne pour le scrutin du 9 juin, fustigeant ceux qui dénoncent l'Union européenne tout en engrangeant les « dividendes silencieux ».						
Vous pouvez partager un article en cliquant sur les icônes de partage en haut à droite de l'article.						
岸田首相、韓国・尹大統領に日米首脳会談の成果を説明　日韓連携も確認						
岸田文雄首相は17日、韓国の尹錫悦（ユンソンヒル）大統領と電話で協議した。首相によると、10日にシンジョンで行われた日米首脳会談での合意事項などについて説明したという。						
協議後、首相は首相官邸で記者団に「日韓、日米との連携をより深化させ、日韓首脳間を含め緊密に意思疎通を図つ	Fumio Kishida, Prime Minister; Yoon Suk-yeol, President of South Korea	岸田首相は17日、韓国の尹大統領と電話で協議し、日米首脳会談の成果を説明、日韓連携も確認した。	ja	no	Prime Minister Kishida held talks by phone with South Korean President Yoon on the 17th, explaining the results of the Japan-U.S. summit and confirming cooperation between Japan and South Korea.	other

B	C	E	F	G	H	I
content	sources	summary	language code	is it English?	English translation	article category
Ante inversores en Washington, Caputo defendió la estrategia fiscal y reafirmó el compromiso para levantar el cepo El jefe del Palacio de Hacienda participó de un seminario organizado por el banco J.P. Morgan WASHINGTON.- El ministro de Economía, Luis Caputo, llegó al seminario del banco de inversión J.P. Morgan en el Park Hyatt hotel poco después de las tres de la tarde, horas después de abandonar Washington.	Luis Caputo, Ministro de Economía Pablo Quirno, Secretario de Finanzas Javier Milei, gobierno Mauricio Macri, gobierno Nicolás Posse, Jefe de Gabinete Alberto Fernández, gobierno Pierre-Olivier Gourinchas, Economista Jefe del FMI	The economy minister defended the government's fiscal adjustment strategy and reiterated the intention to lift currency controls, but said they cannot do so with the current level of reserves unless they receive a fresh injection of money.	en	yes	The economy minister defended the government's fiscal adjustment strategy and reiterated the intention to lift currency controls, but said they cannot do so with the current level of reserves unless they receive a fresh injection of money.	international
Elections européennes : Emmanuel Macron dénonce « l'hypocrisie » du RN Le chef de l'Etat a déploré mercredi depuis Bruxelles que les « peurs profitent aux réponses les plus simplistes » dans la campagne pour le scrutin du 9 juin, fustigeant ceux qui dénoncent l'Union européenne tout en engrangeant les « dividendes silencieux ». Vous pouvez partager un article en cliquant sur les icônes de partage en haut à droite de l'article.	Emmanuel Macron, President of France Valérie Heyer, Head of list for Macron's camp in European elections Jordan Bardella, Head of list for National Rally (Rassemblement National, RN) Raphaël Glucksmann, Head of list for Socialist Party (Parti Socialiste) and Place Publique Gabriel Attal, Prime Minister of France	Emmanuel Macron a déploré que les "peurs profitent aux réponses les plus simplistes", mercredi 17 avril.	fr	no	Emmanuel Macron lamented that "fears profit from the simplest answers", Wednesday April 17.	politics
岸田首相、韓国・尹大統領に日米首脳会談の成果を説明　日韓連携も確認 岸田文雄首相は17日、韓国の尹錫悦（ユンソンニョル）大統領と電話で協議した。首相によると、10日にワシントンで行われた日米首脳会談での合意事項などについて説明したという。 協議後、首相は首相官邸で記者団に「日韓、日米との連携をより深化させ、日韓首脳間を含め緊密に意思疎通を図つ	Fumio Kishida, Prime Minister; Yoon Suk-yeol, President of South Korea	岸田首相は17日、韓国の尹大統領と電話で協議し、日米首脳会談の成果を説明、日韓連携も確認した。	ja	no	Prime Minister Kishida had talks by phone with South Korean President Yoon on the 17th, explaining the results of the Japan-U.S. summit and confirming cooperation between Japan and South Korea.	other



Google Sheets add-on - Anth

docs.anthropic.com/en/docs/build-with-claude/claude-for-sheets#example-setting-parameters

Search... ⌘K

Research News Go to claude.ai

ANTHROPIC

English

Welcome User Guides API Reference Prompt Library Release Notes Developer Newsletter

Use cases >

Models

Security and compliance >

Build with Claude

Define success criteria

Develop test cases

Prompt engineering >

Text generation

Embeddings

Google Sheets add-on

Vision

Tool use (function calling)

Model Context Protocol (MCP)

Computer use (beta)

Prompt caching

Batch processing

PDF support

Citations

Token counting

Multilingual support

Claude for Sheets usage examples

Prompt engineering interactive

Our in-depth prompt engineering interactive guide will walk you through the process of creating effective prompts for your AI models. Whether you're new to prompt engineering or looking to refine your skills, this guide is a great place to start.

Just as with any instance of Claude, you can use the interactive interface to experiment with different prompt variations and see how they affect the model's responses. You can also use the interface to generate code snippets for your own projects.

Prompt engineering workflow

Our Claude for Sheets prompting guide provides a step-by-step workflow for creating effective prompts. It includes tips for structuring prompts, choosing the right tone and style, and testing different variations to find the best results.

Claude for Sheets work

Make a copy of our Claude for Sheets guide to learn more about how to use Claude for Sheets.

The argument-value parameters you might care about most are:

Argument	Description
max_tokens	The total number of tokens the model outputs before it is forced to stop. For yes/no or multiple choice answers, you may want the value to be 1-3.
temperature	The amount of randomness injected into results. For multiple-choice or analytical tasks, you'll want it close to 0. For idea generation, you'll want it set to 1.
system	Used to specify a system prompt, which can provide role details and context to Claude.
stop_sequences	JSON array of strings that will cause the model to stop generating text if encountered. Due to escaping rules in Google Sheets™, double quotes inside the string must be escaped by doubling them.
api_key	Used to specify a particular API key with which to call Claude.

Troubleshooting

NAME2 Error: Unknown function: 'claude'

click me

Claude for Sheets examples (Economist)

File Edit View Insert Format Data Tools Extensions Help

docs.google.com/spreadsheets/d/1YcAOKI7LDZml8ycr2eq_D-u_8jkYXG_BMKrWW8z_ofQ/edit?gid=0#gid=0

E10 | fx

A B C D E F G H I J K L M

1

2 sky

3

4

5 You need to make a copy of this spreadsheet! (File > Make a Copy...)

6 You need to install Claude for Sheets! https://workspace.google.com/marketplace/app/clause_for_sheets/909417792257

7 You will need an API key! Visit <https://jonathansoma.com/it-keys.txt> You add it through Extensions > Claude for Sheets > Add Anthropic API key

8 You might need to go to Extensions > Add Ons > Manage extensions, click the ... next to Claude for Sheets and say "Use in this document"

SEE THE TABS BELOW

home basics food articles (multi-language) comments

Summarize this table

Relaunch to update

J

JS

List the named entities in the passage below. Use a comma to separate the entity and the type of entity. Valid entity types are PERSON, EVENT, ORGANIZATION and LOCATION.

Text: Hi, I'm Jonathan Soma, and I'm giving this talk at the Abraji Conference in São Paulo.



Jonathan Soma - PERSON

Abraji Conference - EVENT

São Paulo - LOCATION

Validating or
constraining
structured output

programming time!

are you using a
remote LLM?

Yes, I'm using
an API

Structured outputs

No, I'm running locally

Is it OpenAI?

Yes, I'm
using GPT

Outlines
(context-free grammars,
constrained generation)

No, Anthropic or
Groq or whatever

Instructor

Pydantic



```
from typing import List, Literal
from pydantic import BaseModel

class Entity(BaseModel):
    text: str
    type: Literal["PERSON", "EVENT", "ORGANIZATION", "LOCATION"]

class EntityList(BaseModel):
    entities: List[Entity]
```

List the named entities in the passage below. Use a comma to separate the entity and the type of entity. Valid entity types are PERSON, EVENT, ORGANIZATION and LOCATION.

Text: Hi, I'm Jonathan Soma, and I'm giving this talk at the Abraji Conference in São Paulo.

Pydantic



```
from pydantic import BaseModel, Field
from typing import Literal

class Comment(BaseModel):
    name: str = Field(..., description="Commenter's name")
    product: str = Field(..., description="Food product")
    food_category: Literal['produce', 'canned goods', 'candy', 'other'] = Field(
        ..., description="Product category"
    )
    alternative_category: str = Field(
        ..., description="Alternative category if 'food_category' is 'other'"
    )
    emotion: Literal['positive', 'negative'] = Field(
        ..., description="Comment sentiment"
    )
```

I am irate about the broccoli incident, I am never shopping at your store again. Take me off of your broccoli mailing list immediately.

Jackary Baloneynoze

jackary.baloneynoze@example.com

Yes, I'm using OpenAI
through an API

Structured outputs



```
client = OpenAI()

class Entity(BaseModel):
    text: str
    type: Literal["PERSON", "EVENT", "ORGANIZATION", "LOCATION"]

class EntityList(BaseModel):
    entities: List[Entity]

completion = client.beta.chat.completions.parse(
    model="gpt-4o",
    messages=[
        {"role": "system", "content": "Extract the entities from the document."},
        {"role": "user", "content": content},
    ],
    response_format=EntityList, ←
)

response = completion.choices[0].message.parsed
```

Only works with OpenAI

Only does this one thing

No, I'm using another
LLM provider

Instructor



```
# Initialize OpenAI client with instructor
client = instructor.patch(OpenAI())

class Entity(BaseModel):
    text: str
    type: Literal["PERSON", "EVENT", "ORGANIZATION", "LOCATION"]

class EntityList(BaseModel):
    entities: List[Entity]

# Perform the request using instructor's parsed response handling
response = client.chat.completions.create(
    model="gpt-4o",
    messages=[
        {"role": "system", "content": "Extract the entities from the document."},
        {"role": "user", "content": content},
    ],
    response_model=EntityList, # Instructor automatically parses response
)
```



Works with any provider

Easy to add validation

Lots of examples and well-documented use cases

Responses are NOT guaranteed to be correct

No, I'm running my
model locally

Outlines



```
class Entity(BaseModel):
    text: str
    type: Literal["PERSON", "EVENT", "ORGANIZATION", "LOCATION"]
```

```
class EntityList(BaseModel):
    entities: List[Entity]
```

```
# Our local model
model = outlines.models.transformers("microsoft/Phi-3-mini-4k-instruct")
```

```
# Construct structured sequence generator
generator = outlines.generate.json(model, EntityList)
response = generator(f"""
    ## INSTRUCTIONS
```

```
    Return the information below in the following JSON schema:
{EntityList.model_json_schema()}

    ## CONTENT
```

```
{content}
""")
```



Only works locally

Examples are not wonderful

Lots of little tweaks can be made for requirements without reasks or retries

Speeds up inference

Can be hooked into any type of model (local multimodal!)

Generative AI chatbots
are AI, but AI isn't
just chatbots.

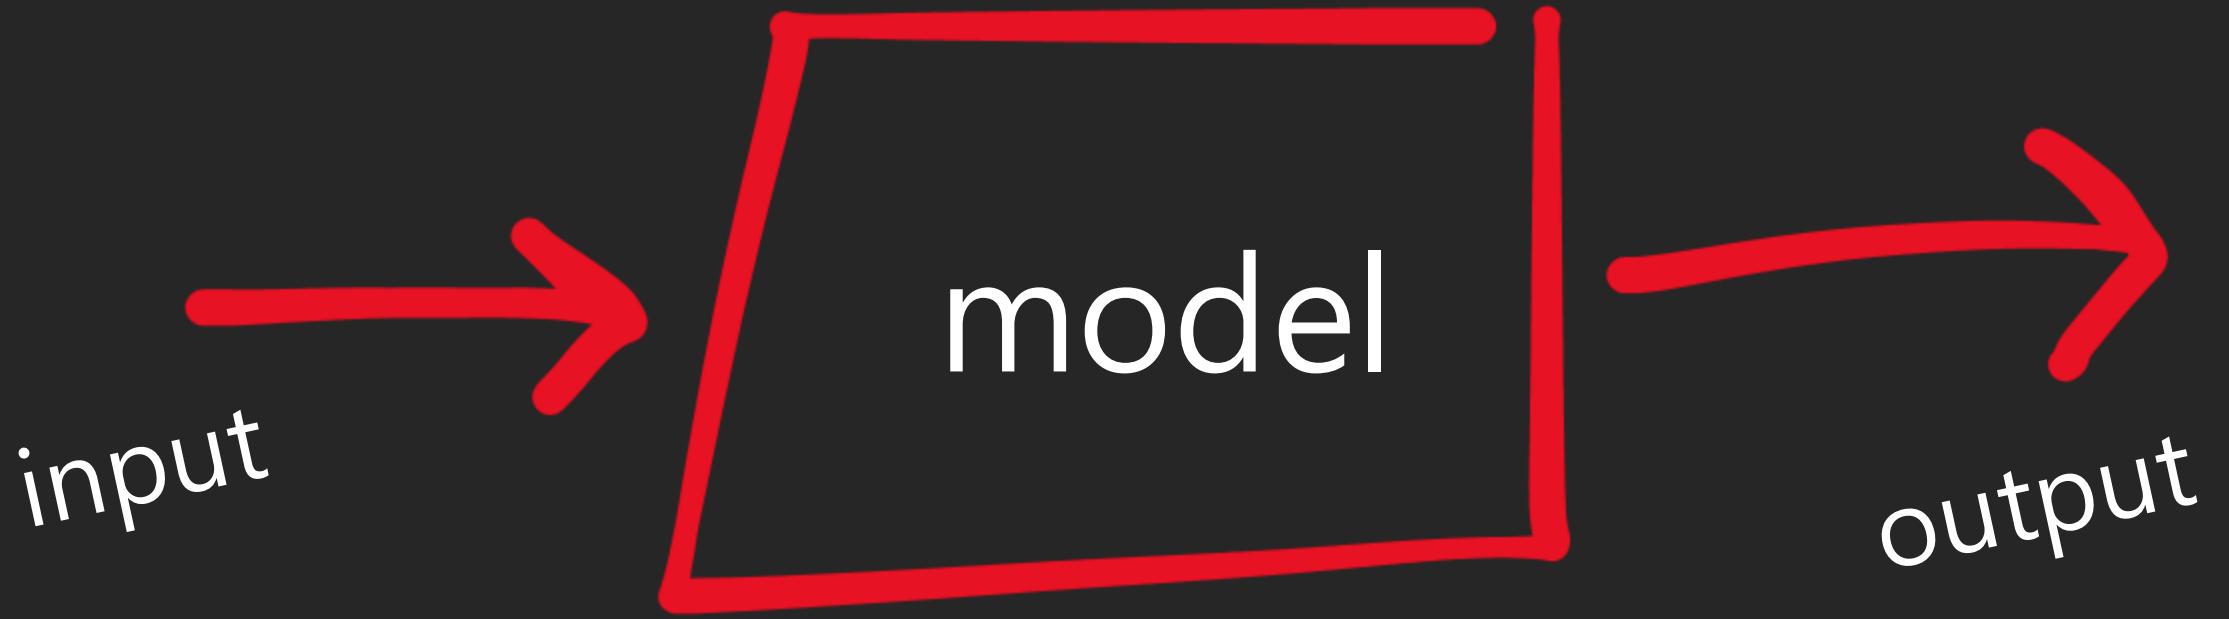
AI has a discoverability problem

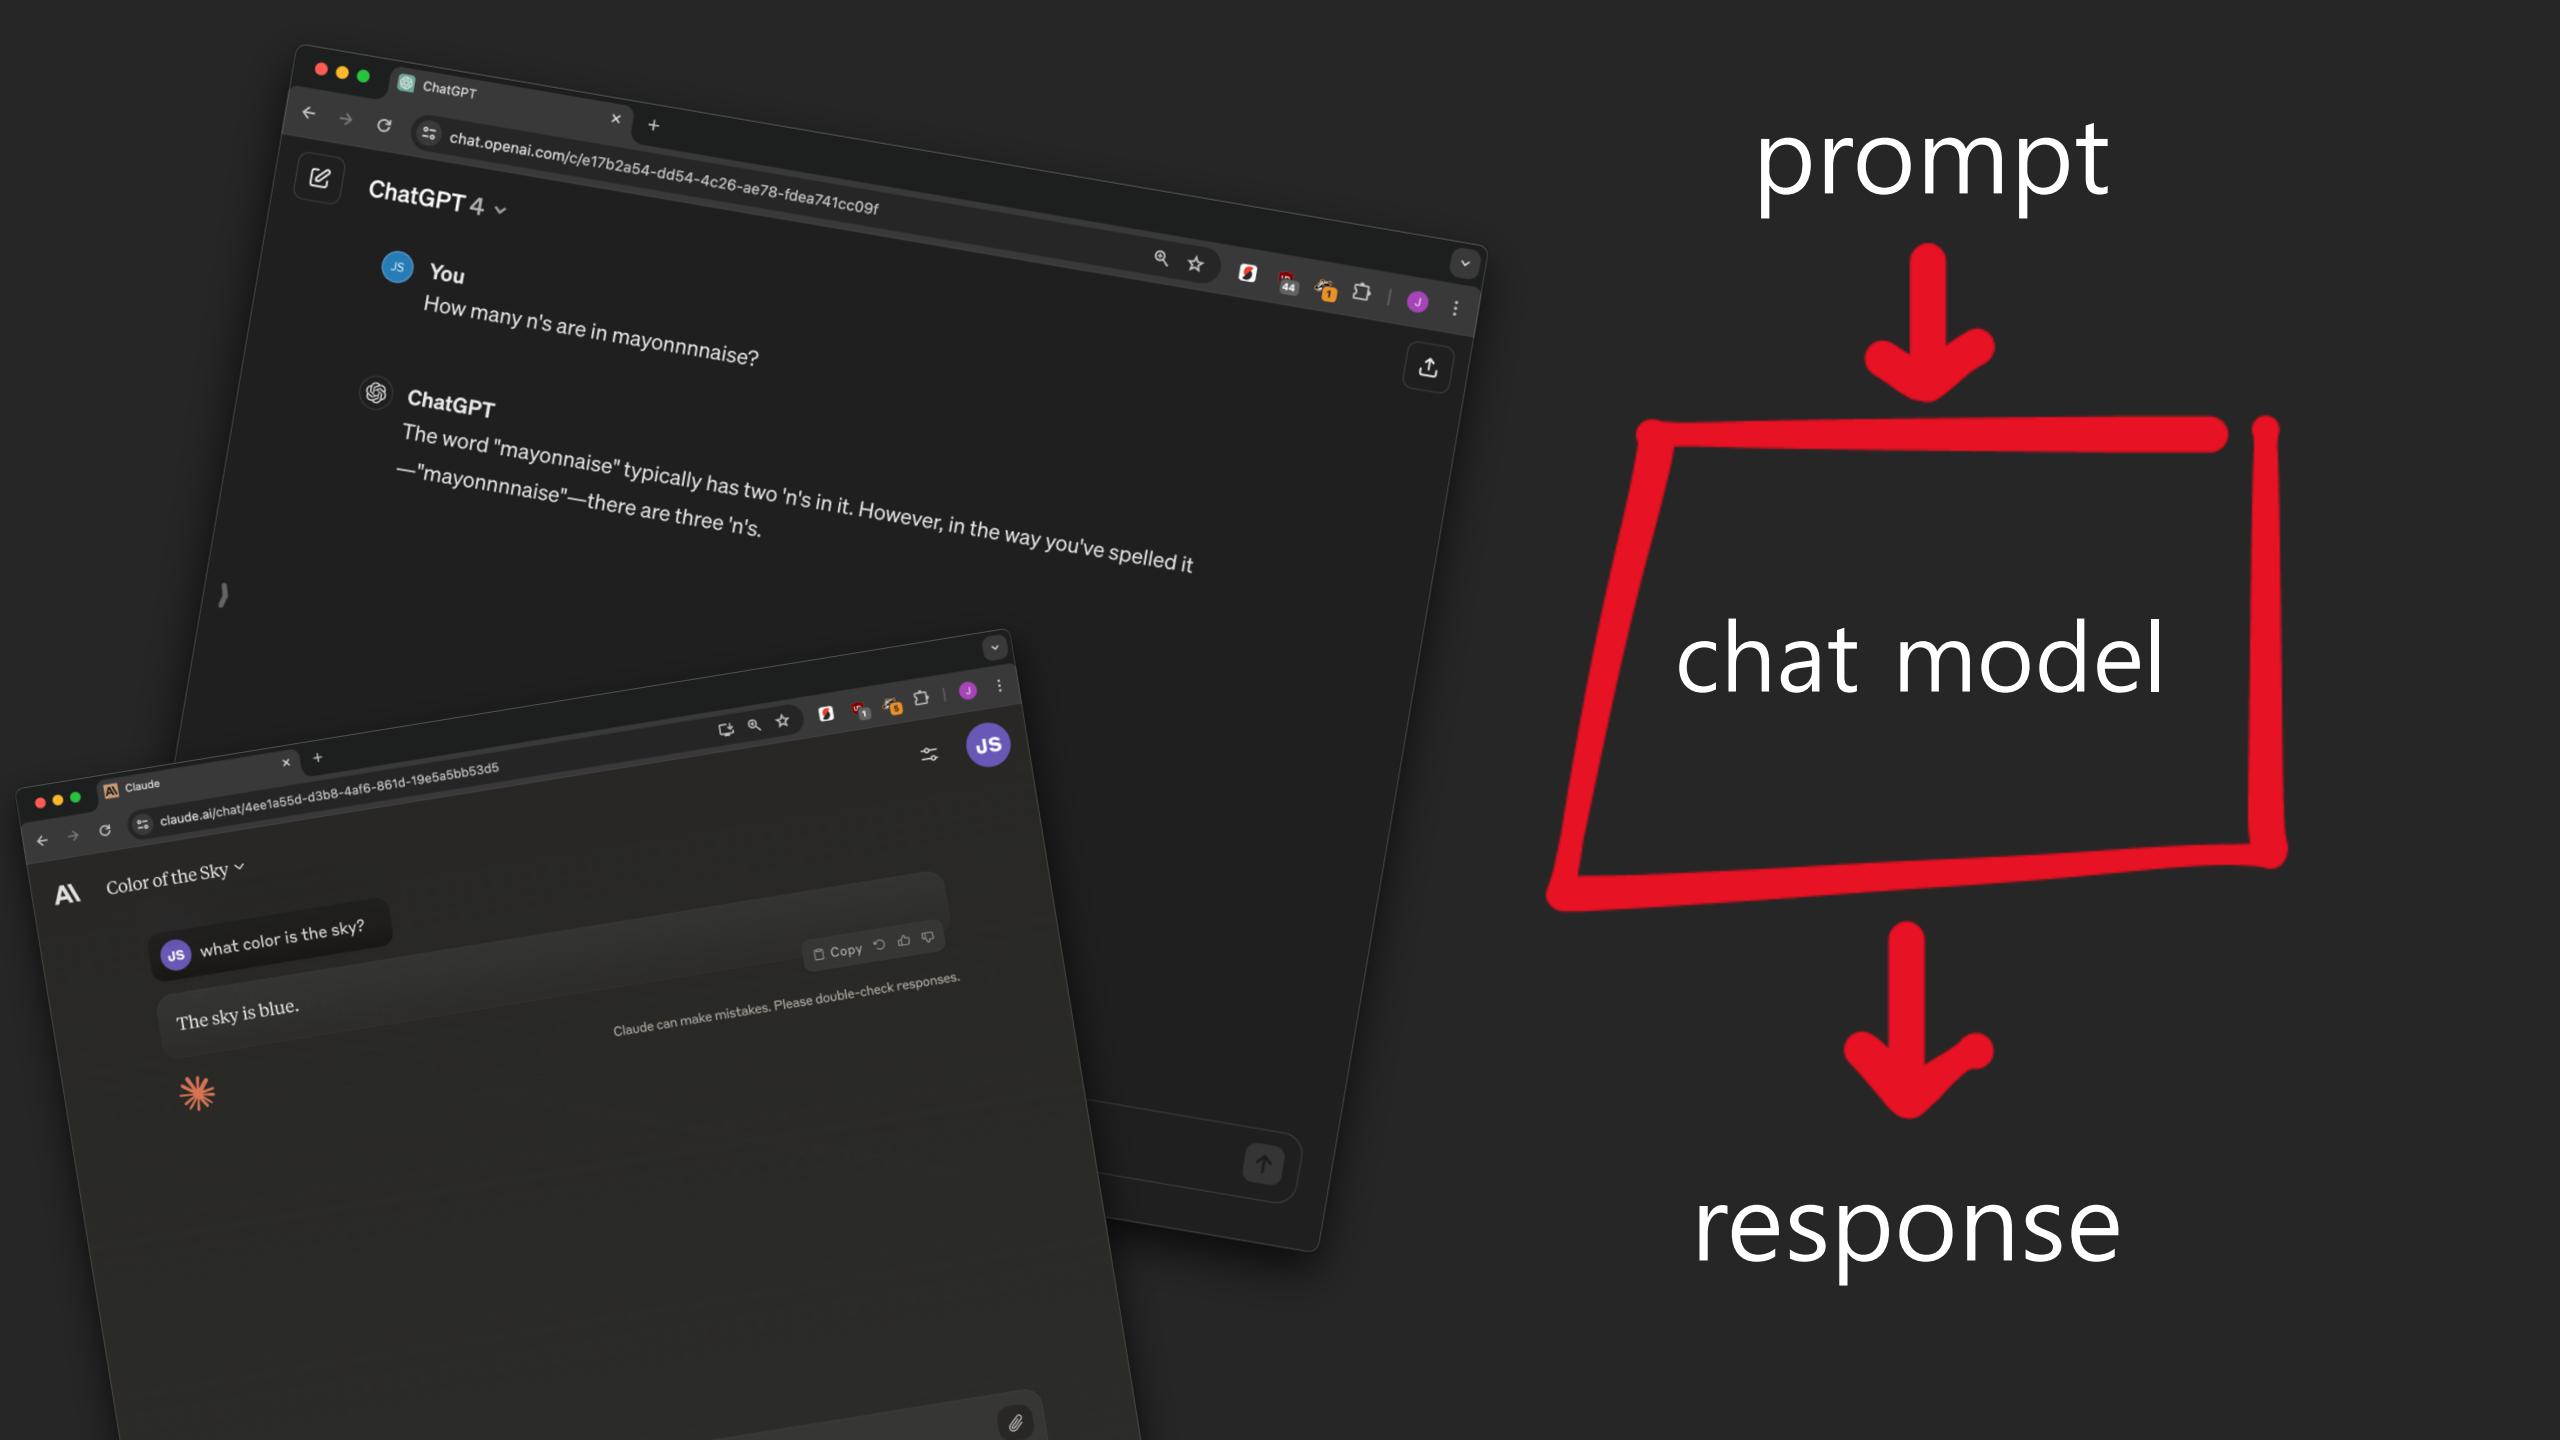


models

"AI"







prompt

chat model

response

ChatGPT 4

You

How many n's are in mayonnnnnaise?

ChatGPT

The word "mayonnaise" typically has two 'n's in it. However, in the way you've spelled it—"mayonnnnnaise"—there are three 'n's.

AI Color of the Sky

what color is the sky?

The sky is blue.

Claude can make mistakes. Please double-check responses.



Tasks Libraries Datasets Languages Licenses Other

Filter Tasks by name

Multimodal

Audio-Text-to-Text Image-Text-to-Text
Visual Question Answering Document Question Answering Video-Text-to-Text
Any-to-Any

Computer Vision

Depth Estimation Image Classification
Object Detection Image Segmentation
Text-to-Image Image-to-Text Image-to-Image
Image-to-Video Unconditional Image Generation
Video Classification Text-to-Video
Shot Image Classification Mask Generation

Hugging Face

Text Classification Token Classification
Table Question Answering Question Answering
Zero-Shot Classification Translation
Summarization Feature Extraction
Text Generation Text2Text Generation
Fill-Mask Sentence Similarity

Models 1,344,188

Filter by name

Full-text search

↑ Sort: Trending

deepseek-ai/DeepSeek-R1
Text Generation • Updated 2 days ago • ↓ 189k • 4.17k

deepseek-ai/Janus-Pro-7B
Any-to-Any • Updated about 24 hours ago • ↓ 19.6k • 1.47k

deepseek-ai/DeepSeek-V3
Text Generation • Updated 4 days ago • ↓ 325k •

deepseek-ai/DeepSeek-R1-Distill-Qwen-7B
Text Generation • Updated 2 days ago • ↓ 132k •

deepseek-ai/DeepSeek-R1-Distill-Qwen-1.5B
Text Generation • Updated 2 days ago • ↓ 150k • 498

deepseek-ai/DeepSeek-R1-Zero
Text Generation • Updated 2 days ago • ↓ 12.3k • 559

deepseek-ai/DeepSeek-R1-Distill-Llama-70B
Text Generation • Updated 2 days ago • ↓ 47.4k • 326

HKUSTAudio/Llasa-3B
Text-to-Speech • Updated 2 days ago • ↓ 3.33k • 360

deepseek-ai/Janus-Pro-1B
Any-to-Any • Updated about 24 hours ago • ↓ 3.11k • 209

deepseek-ai/DeepSeek-R1-Distill-Llama-8B
Text Generation • Updated 2 days ago • ↓ 57.7k • 287

openbmb/MiniCPM-o-2_6
Any-to-Any • Updated 2 days ago • ↓ 86.3k • 856

cyberagent/DeepSeek-R1-Distill-Qwen-32B-Japanese
Text Generation • Updated 1 day ago • ↓ 1.39k • 148

ostris/Flex.1-alpha
Text-to-Image • Updated 10 days ago • ↓ 13.1k • 305

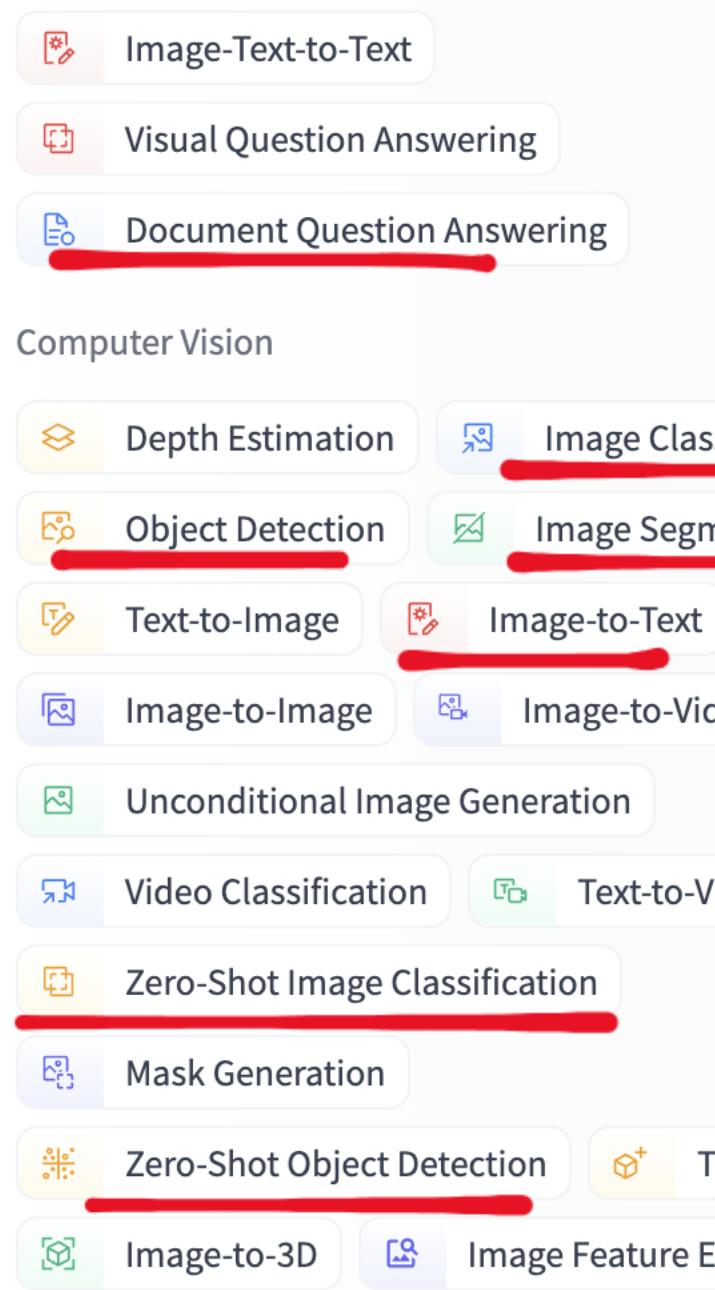
deepseek-ai/DeepSeek-R1-Distill-Qwen-14B
Text Generation • Updated 2 days ago • ↓ 44k • 230

deepseek-ai/DeepSeek-R1-Distill-Qwen-7B
Text Generation • Updated 2 days ago • ↓ 72.5k • 237

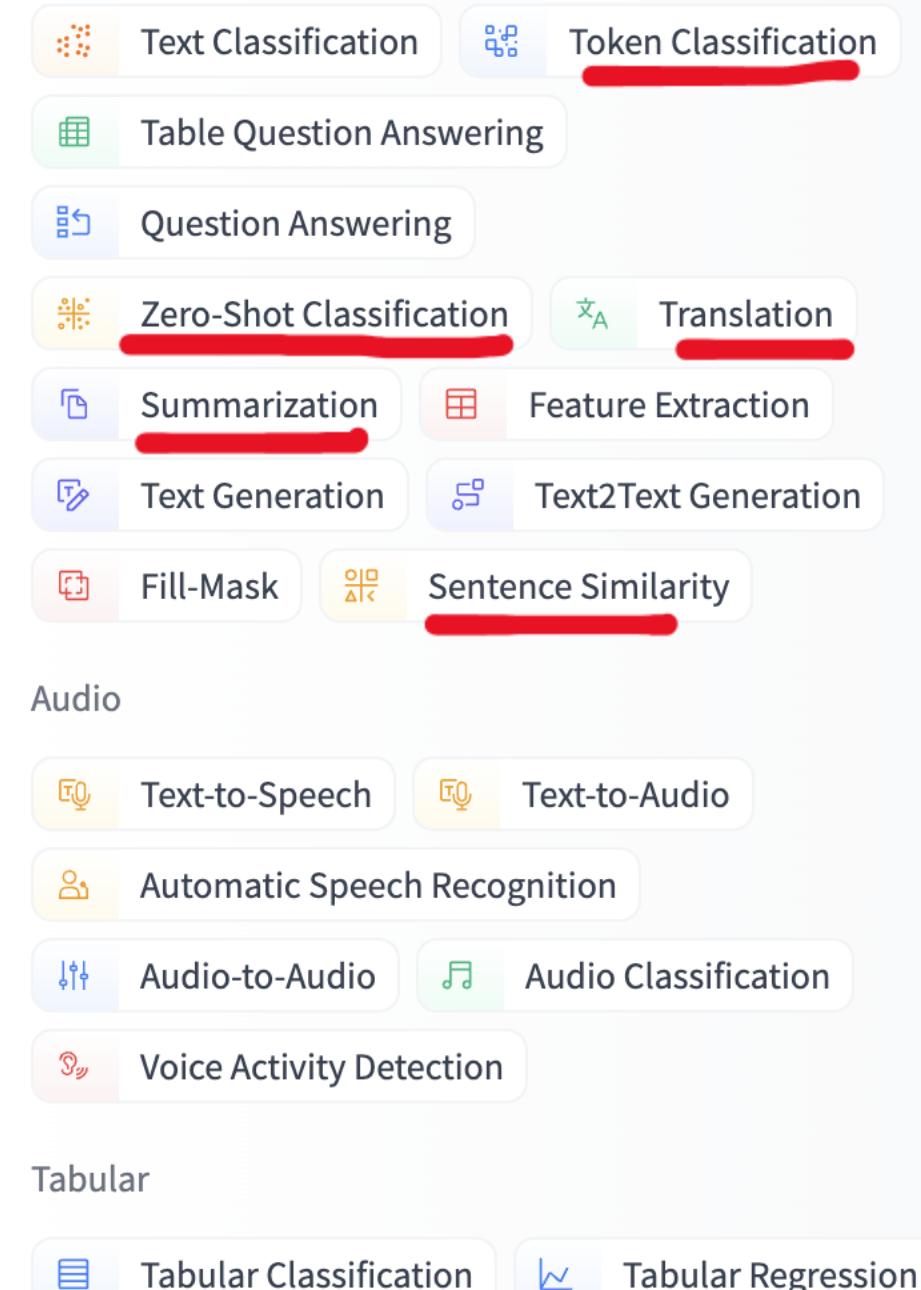
Qwen/Qwen2.5-14B-Instruct-1M
Text Generation • Updated about 8 hours ago • ↓ 1.8k • 138

Models 1,344,188

Multimodal



Natural Language Processing



document

fx =CLAUDEEXTRACT("Is this piece of legislation about retirement? Answer YES or NO", A2)

A

Louisiana state legislation (completed in class) Saved to Drive

File Edit View Insert Format Data Tools Extensions Help

G3 | fx

	A	E	F	G	H
1	bill text	about retirement?	about employment?		
1	> HLS 24RS-94 **[ENGROSSED]{.underline}** > > 2024 Regular Session > > HOUSE BILL NO. 8 > > BY REPRESENTATIVES MELERINE AND HUGHES > > SCHOOLS/HIGH SCHOOL: Prohibits the use of an appeals pro > certain students >	no	no		
2	> HLS 24RS-88 **[REENGROSSED]{.underline}** > > 2024 Regular Session > > HOUSE BILL NO. 9 > > BY REPRESENTATIVE ILLG > > RETIREMENT/DISTRICT ATTY: Provides for survivor benefits o > the >	yes	no		
3					

+ Sheet1 Sheet3

classification
model

category

text



named entity
recognition
model



entities

List the named entities in the passage below. Use a comma to separate the entity and the type of entity. Valid entity types are PERSON, EVENT, ORGANIZATION and LOCATION.

Text: Hi, I'm Jonathan Soma, and I'm giving this talk at the Abraji Conference in São Paulo.



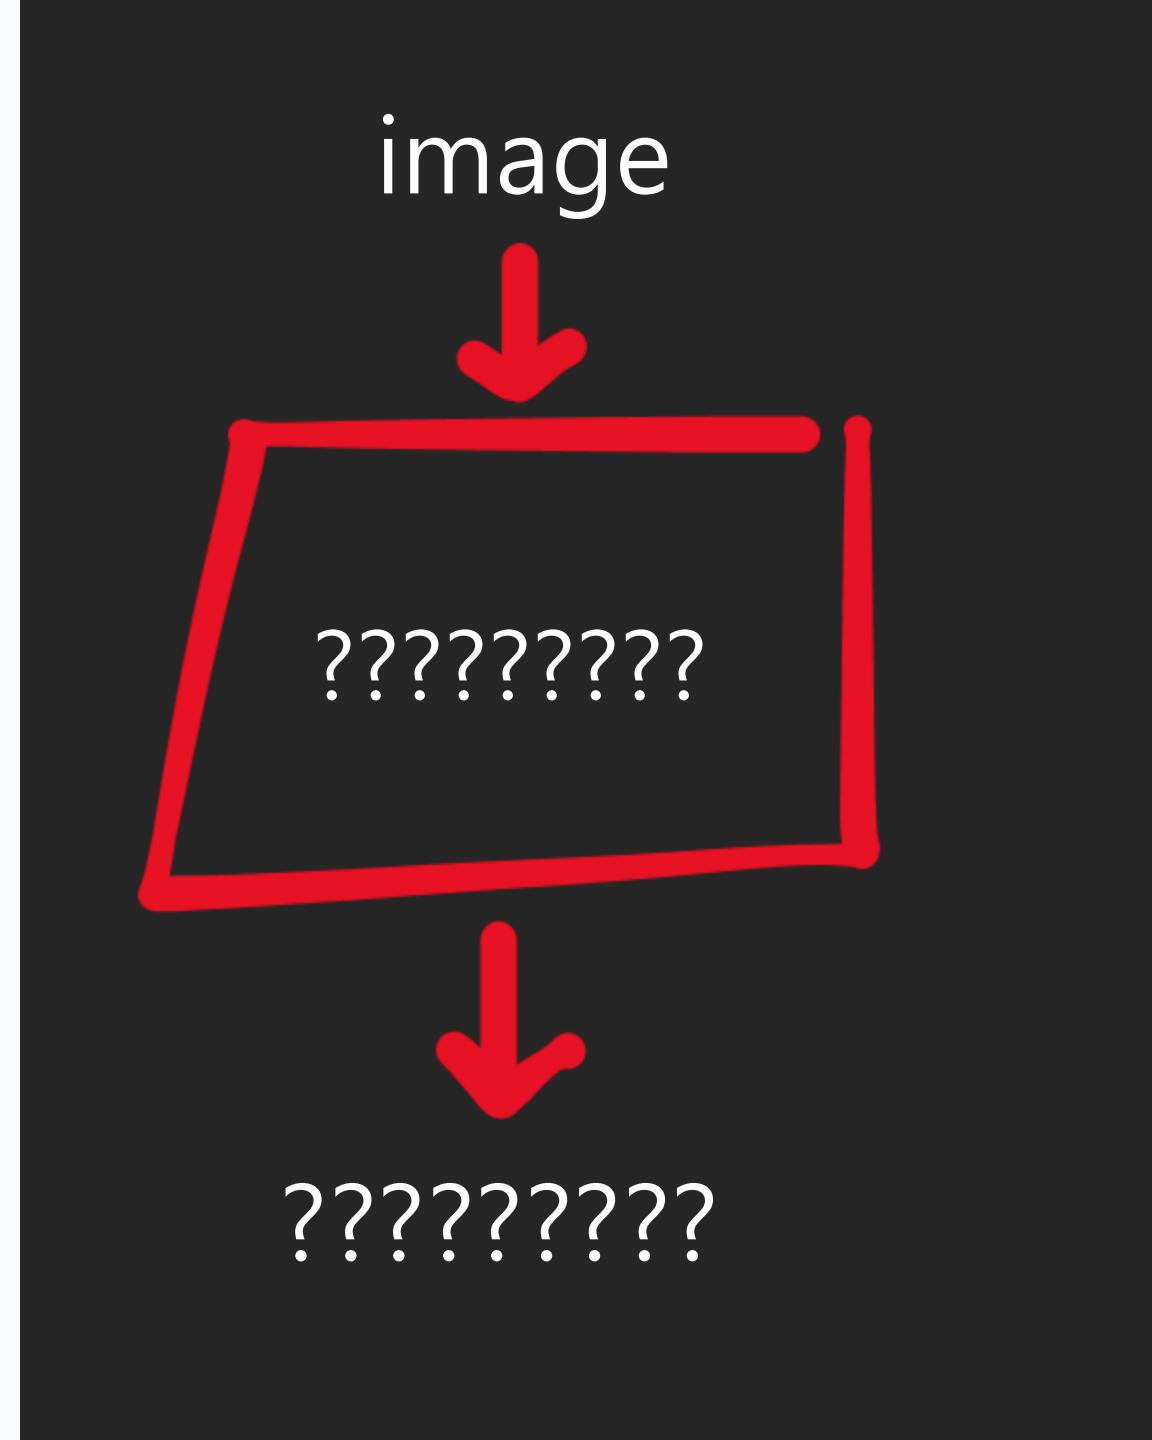
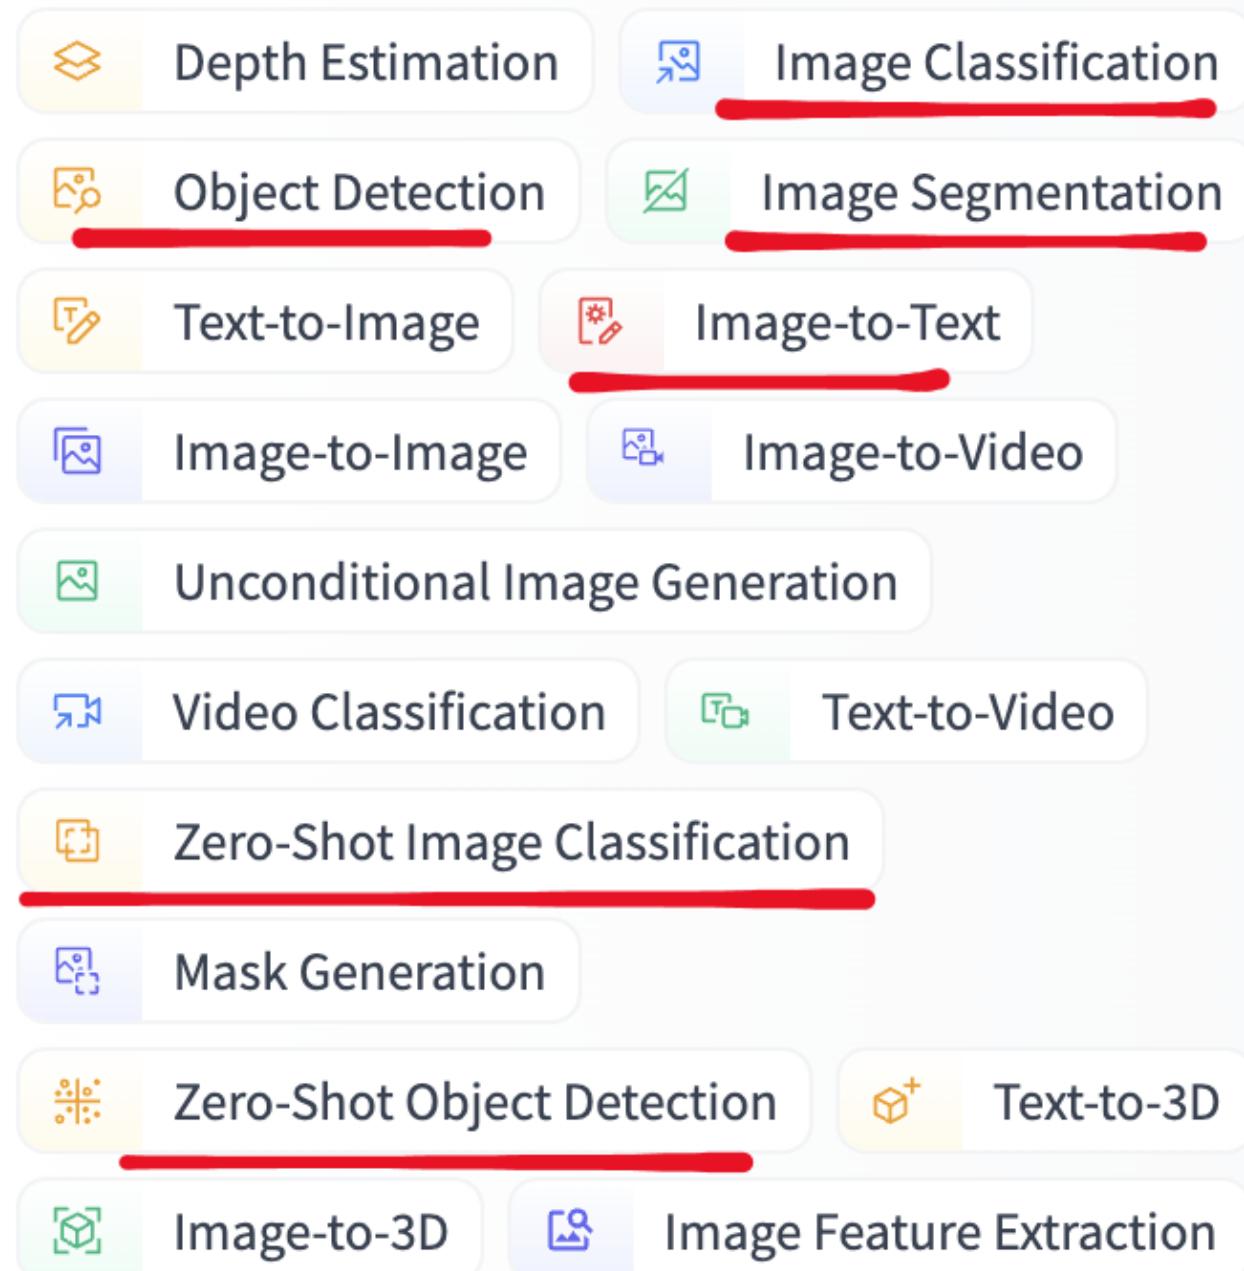
Jonathan Soma - PERSON
Abraji Conference - EVENT
São Paulo - LOCATION

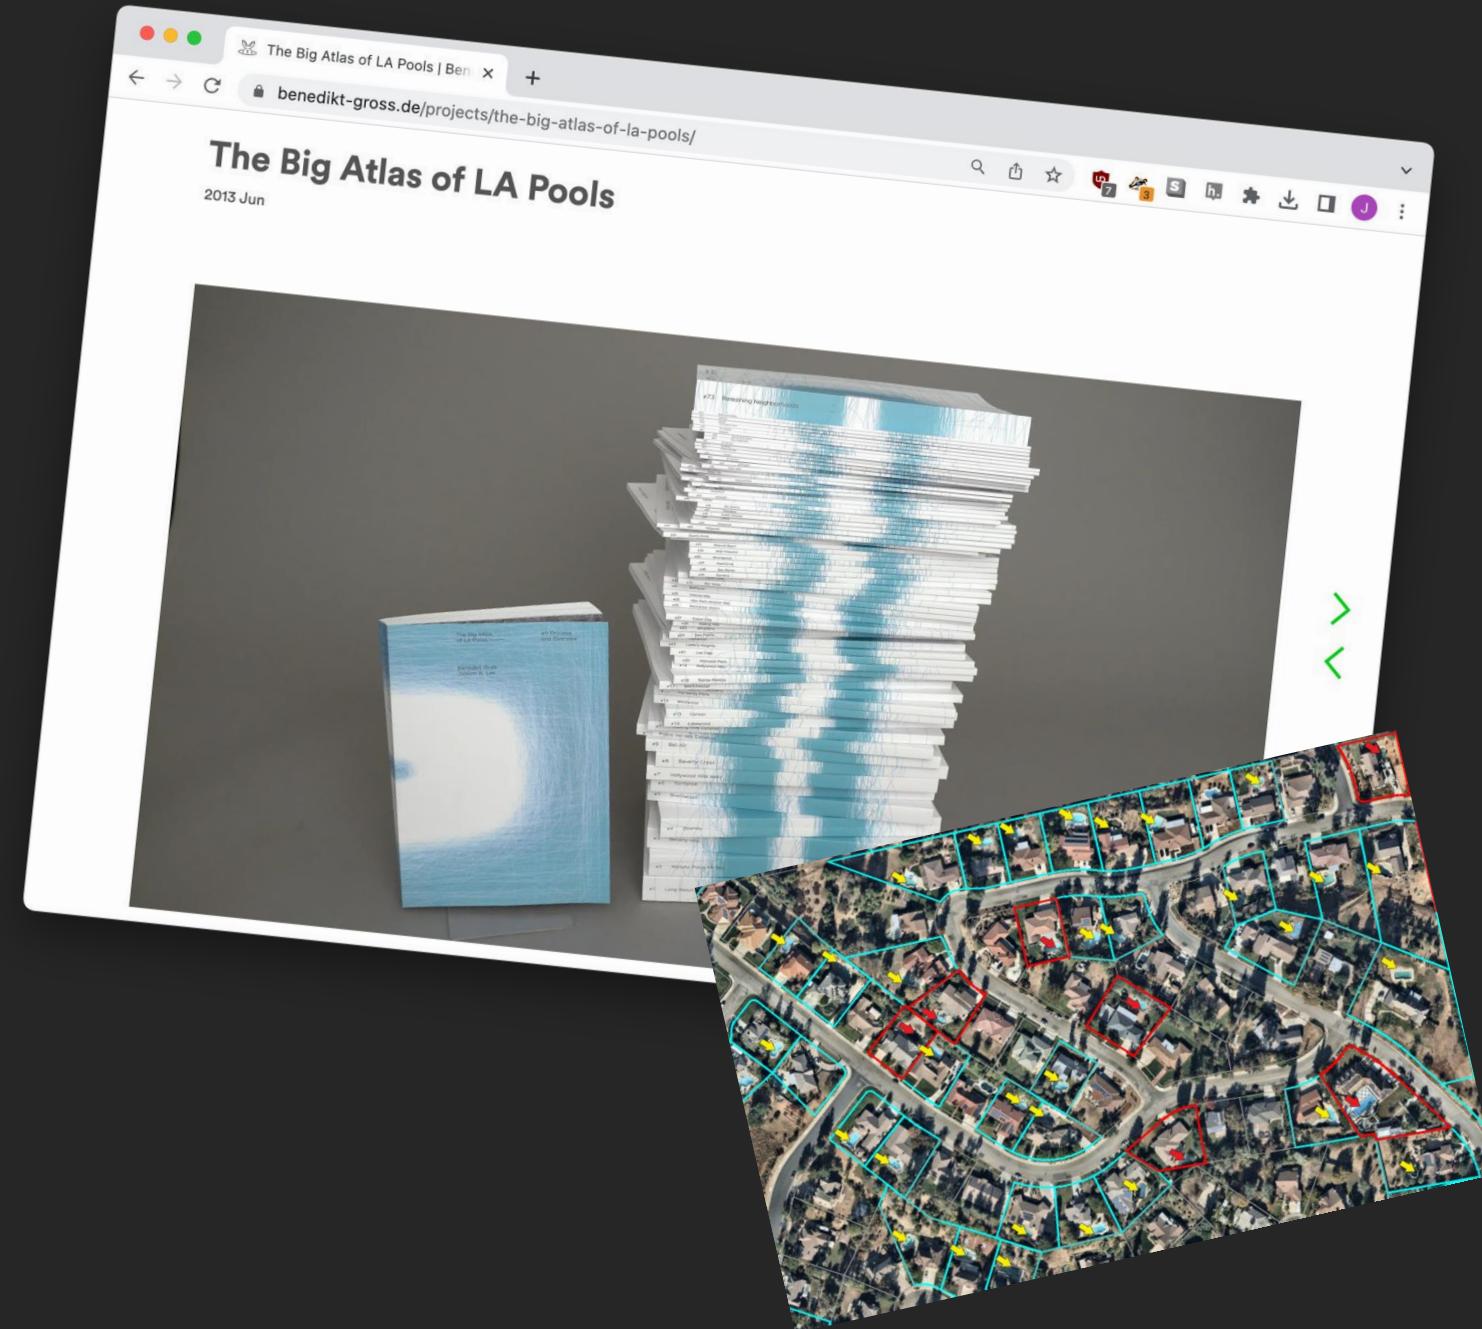
Chatbots can almost
do too much.

knowing the original tasks helps
you understand what's possible



Computer Vision





image



object
detection
model



swimming
pools

TK Leprosy of the land

texty.org.ua/d/2018/amber_eng/ Relaunch to update

Leprosy of the land



THOUSANDS HECTARES OF LAND TURNED INTO LUNAR LANDSCAPES DUE TO ILLEGAL AMBER MINING. OUR MACHINE MODEL FOUND TRACES OF SUCH ACTIVITY ON SATELLITE IMAGES FOR THE AREA OF 70,000 SQUARE KM.

Illegal extraction of gemstones provide "jobs" to thousands of people who live near such deposits in Volyn, Zhytomyr, and most of all, in Rivne region of a country.

A few decades ago one could hardly imagine how to make profit out of the "solar stone". In the past people could dig up amber in their backyards and burn it in stoves as



Ukraine has a world amber market share of 30-40%, it is

image



object
detection
model



illegal amber
mines

How They Did It: Uncovering a Vast Network of Illegal Mining in Venezuela

by Mariel Lozada • June 2, 2022

In the Venezuelan Amazon, many hours away from the country's capital, traces of the devastation caused by illegal mining can be seen from the sky.

Although this area is officially protected, and home to a number of Indigenous communities, organized crime groups have seized control of illicit quarries, destroying the rainforest, and even built runways to take their bounties overseas, according to NGOs like [Human Rights Watch](#). While observers suspected that the country's devastating economic crisis had made the problem of illegal mining worse, it was hard to grasp the real magnitude of the problem.

But a collaboration between the Spanish newspaper El País and the Global Investigative Journalism Network allowed us

Image: Screenshot

READ THIS ARTICLE IN
বাংলা English Español

SHARE
Tweet Facebook Print LinkedIn WhatsApp Telegram Reddit Republish

STORIES
TOPICS
» How They Did It
» Methodology

Global Investigative Journalism Network

Search

Relaunch to update

Donate

CORREDOR FURTIVO

image

object
detection
model

airstrips

G Investigating stories with Mac

newsinitiative.withgoogle.com/resources/trainings/hands-on-machine-learning/investigating-stories-with-machine... ☆

Relaunch to update

Google News Initiative Resources About

< HANDS-ON MACHINE LEARNING

Lesson 2 of 7 Investigating stories with Machine Learning

Leprosy of the Land, an investigation by Texty

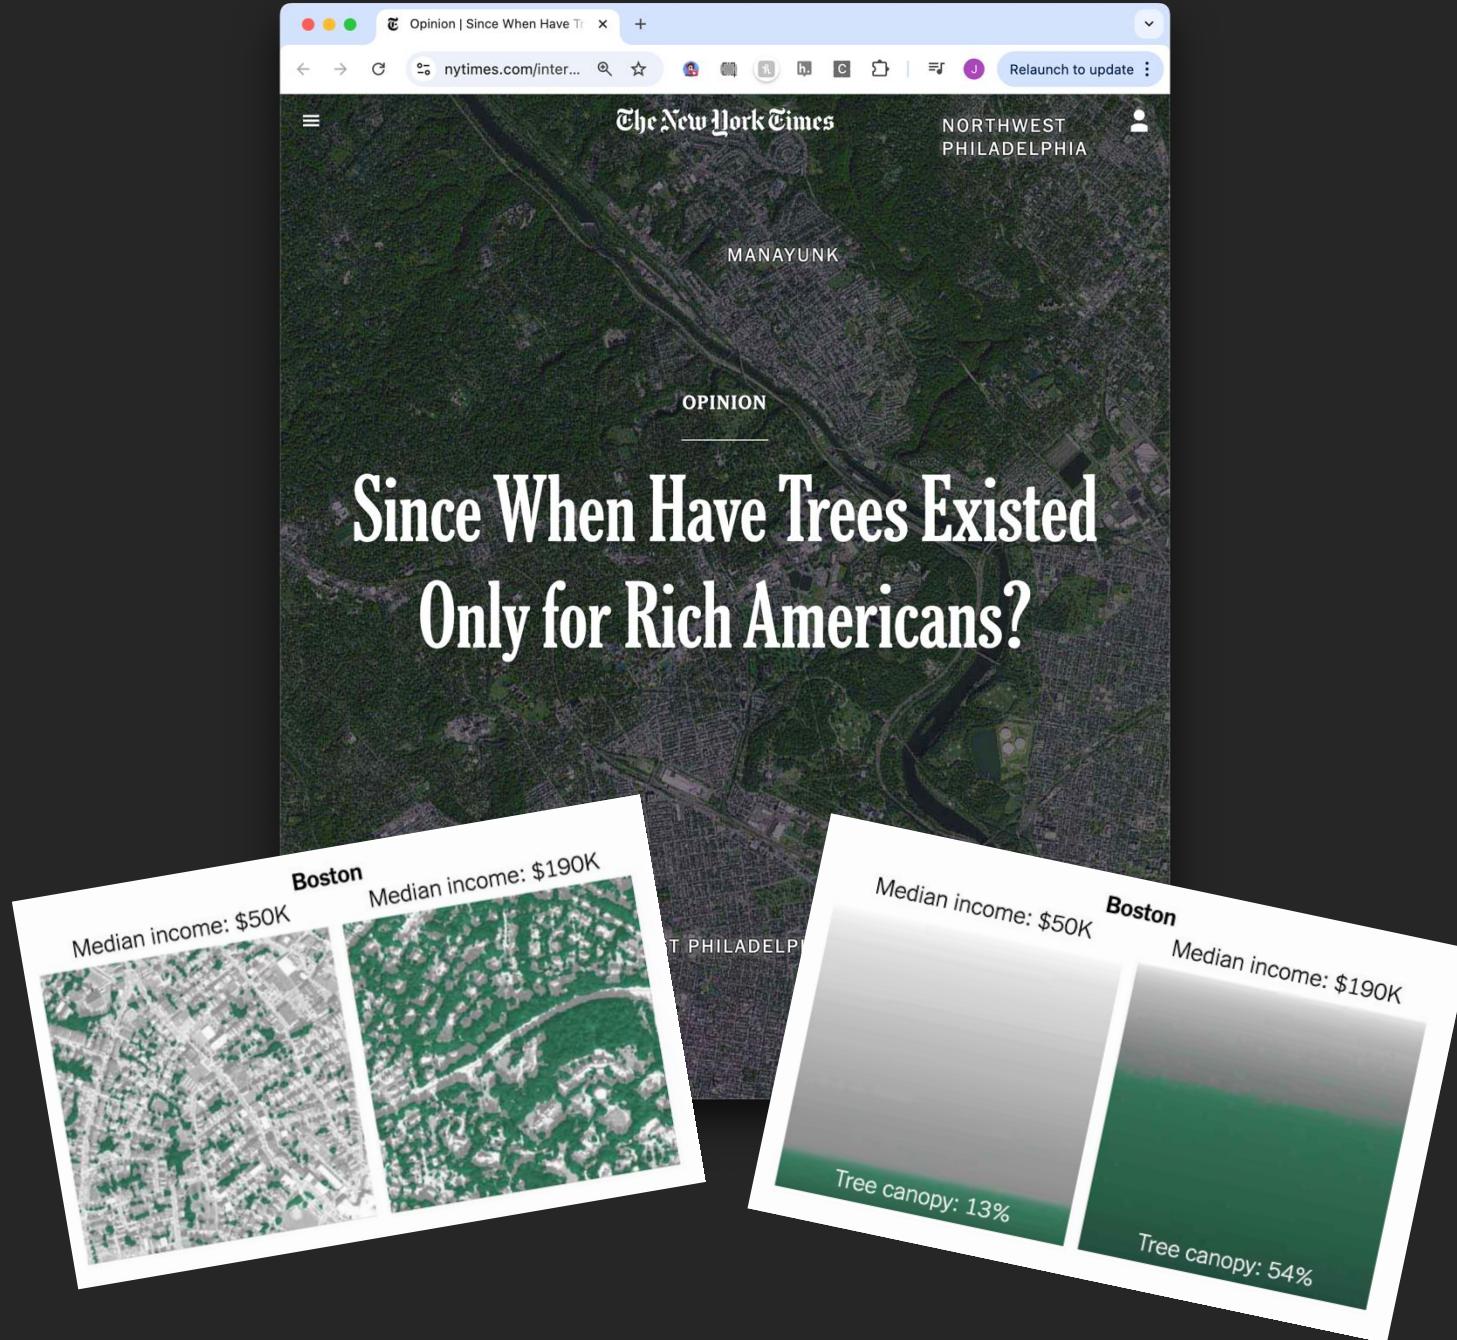
In 2018, Ukrainian data journalism agency Texty published [Leprosy of the Land](#), an investigation in which they used machine learning techniques to detect cases of illegal amber mining across Ukraine.

First, an algorithm divided sections of satellite images into visually uniform subsections. So if an image was half green forest and half dirt field, it would split the image into those two subsections.

Another algorithm found which subsections most resembled the existing examples of amber mining, which have a distinctive pockmark-like pattern of holes in the ground.

Finally, the journalists examined the examples the algorithm found, to make sure that what it thought looked like amber mining wasn't actually something else, like deforestation.

Rivne region Rokytno district, Verchysia
51.5824, 27.1043



image



image
segmentation
model



greenery

Maskformer COCO Vegetation

huggingface.co/spaces/thiagohersan/maskfor...

Spaces thiagohersan/maskformer-coco-vegetation-gradio

Relaunch to update

Maskformer (large-coco)

Using [facebook/maskformer-swin-large-coco](#) model to calculate percentage of pixels in an image that belong to vegetation.

Input Image



Vegetation



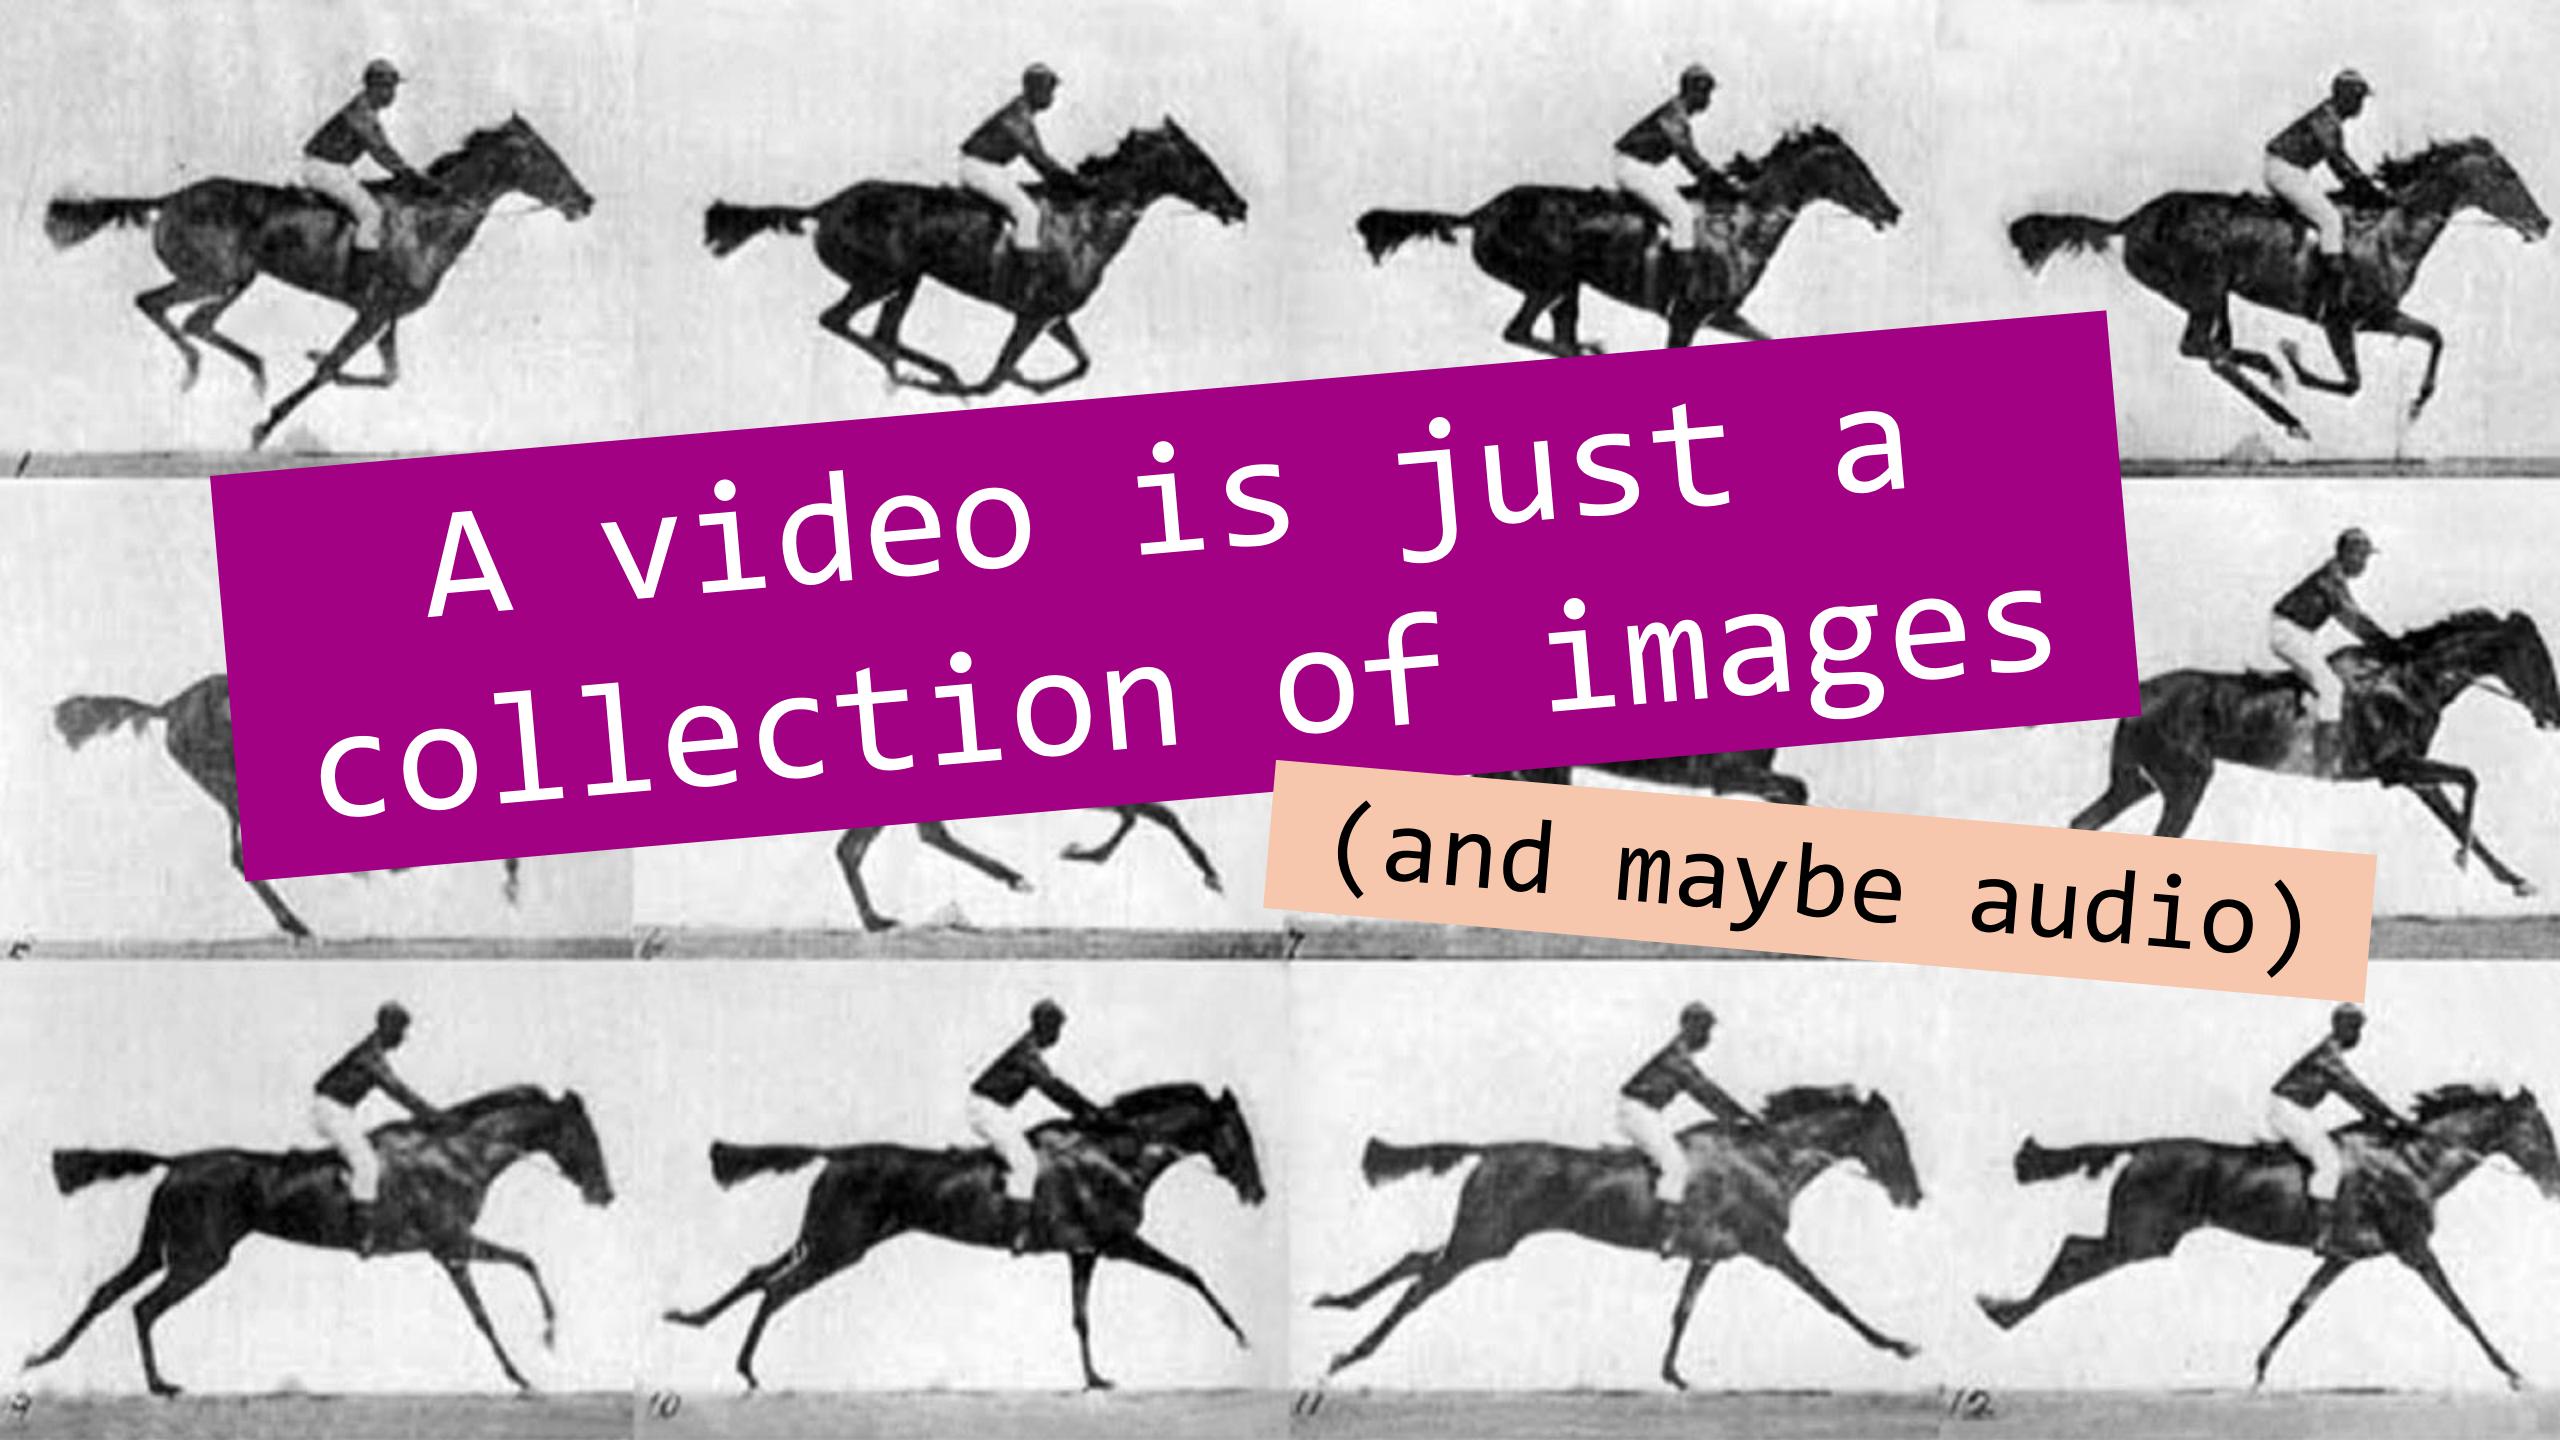
Info

Object Label	Pixel Percent	Square Length
vegetation	40.77 %	0.64 m

Clear Submit

Examples



A sequence of frames from a video showing a horse race. The frames are arranged in a staggered grid, with some frames partially visible at the top and bottom edges. Each frame depicts a jockey on a dark-colored horse in mid-gallop. The background is a plain, light color.

A video is just a
collection of images

(and maybe audio)

NYC Congestion Pricing: How

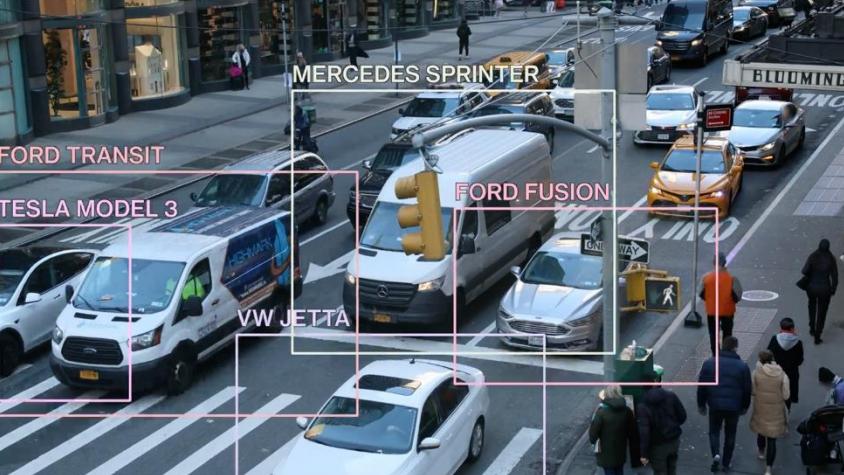
bloomberg.com/grap... Relaunch to update

Bloomberg the Company & Its Products | Bloomberg Terminal Demo Request | Bloomberg Anywhere Remote Login | Bloomberg Customer Support

Bloomberg

Sign In Subscribe

Live TV Markets Economics Industries Tech Politics Businessweek Opinion More US Edition



Vehicles are classified just after they enter the congestion pricing zone in Manhattan.

Is NYC's Congestion Pricing Working? Fewer Private Cars Are On the Road (For Now)

A Bloomberg analysis of 75,000 vehicles found fewer private cars and more cabs in Manhattan's traffic mix.

By Aaron Gordon, Priyanjana Bengani, Jeff Kao, Marie Patino and Jason Kao for Bloomberg CityLab
January 17, 2025

[Facebook](#) [Twitter](#) [LinkedIn](#) [Email](#) [Gift this article](#)

The number of cars entering the zone below 60th Street is down about 8% from baseline traffic levels, according to preliminary data released on Monday by the Metropolitan Transportation Authority. The data also shows traffic speeds into the zone are largely up, a finding independently corroborated by data from [HERE Technologies](#), a digital mapping and location data company.

“So far, so good,” said former New York City deputy traffic commissioner Bruce Schaller. “And about what we expected.”

However, the data still misses a key question: What vehicles disappeared?

Situated on 59th Street, overlooking the congestion pricing zone perimeter, the newsroom aimed a high-definition **camera** at Lexington Avenue – a major thoroughfare into the zone from the Upper East Side – and collected more than a hundred hours of daytime traffic imagery before and after congestion pricing went into effect.



Yellow Taxis
+7 percentage points

Taxi license plates begin with **Y** and end in **C**



For Hire (Includes Uber, Lyft)
-1 pct. pts

For-hire vehicle license plates begin with **T** and end in **C**



Commercial
No change

Includes vans, box trucks and large freight vehicles

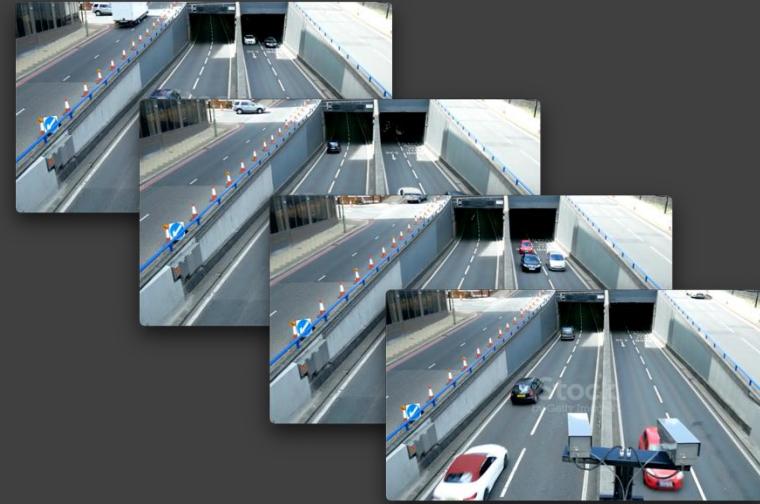


Private vehicles
-6 pct. pts.

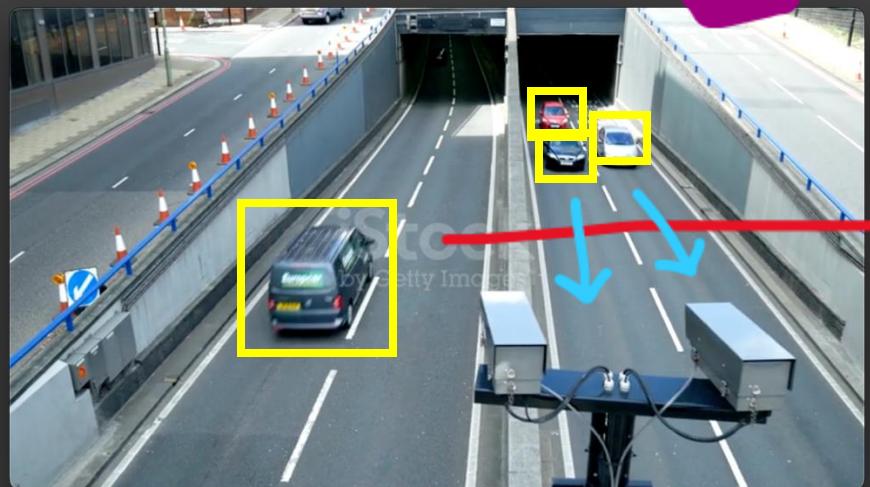
All remaining vehicles were classified as private



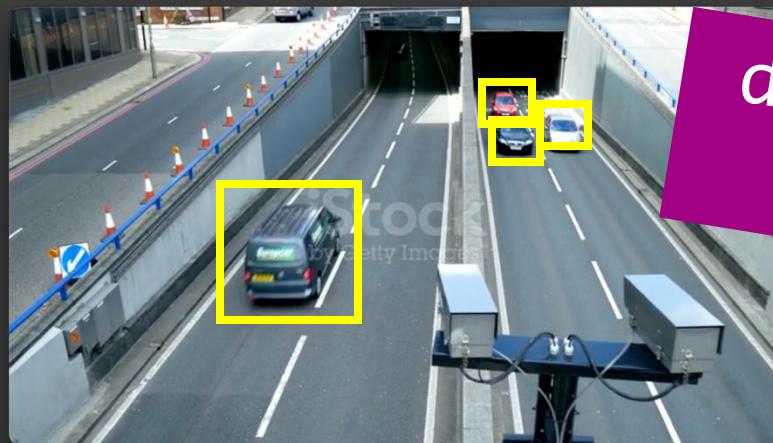
split into
images



track them



count them



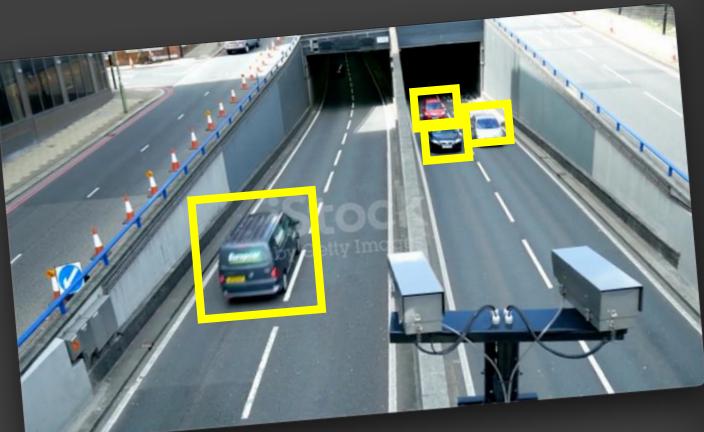
detect your
objects

split into
images

```
● ● ●  
import supervision as sv  
  
video_file = "istockphoto-534232220-640_adpp_is.mp4"  
frame_generator = sv.get_video_frames_generator()  
  
for frame in frame_generator:  
    # You have a frame!
```

detect your
objects

```
● ● ●  
from inference import get_model  
  
model = get_model("yolov8n-640")  
for frame in frame_generator:  
    result = model.infer(frame)[0]  
    detections = sv.Detections.from_inference(result)
```



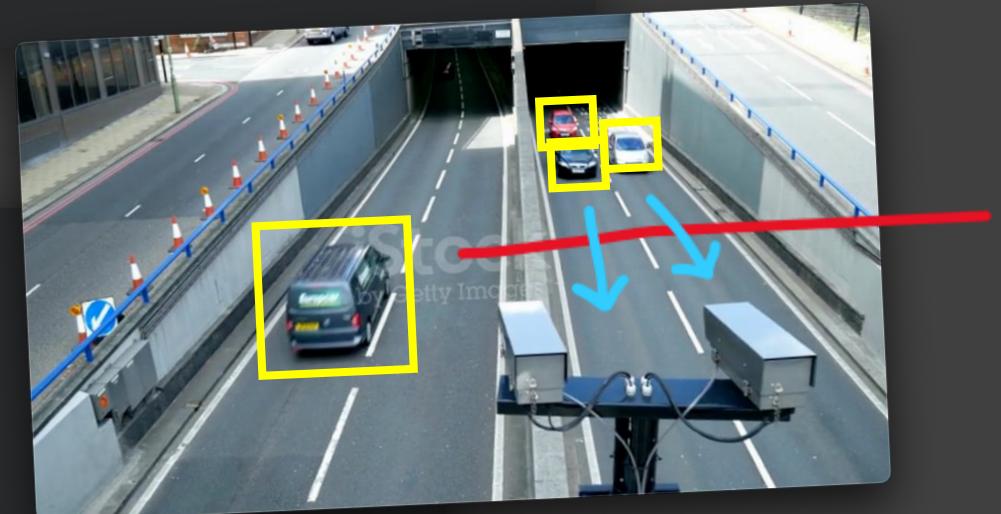
track them



```
byte_track = sv.ByteTrack()  
byte_track.reset()  
  
for frame in frame_generator:  
    # ...  
    detections = byte_track.update_with_detections(detections)  
    # ...
```

count them

```
start = sv.Point(200, 175)  
end = sv.Point(700, 175)  
line_zone = sv.LineZone(start, end)  
  
for frame in frame_generator:  
    # ...  
    line_zone.trigger(detections)  
    # ...
```





annotate
the video

```
box_annotator = sv.BoxAnnotator()
label_annotator = sv.LabelAnnotator()
line_zone_annotator = sv.LineZoneAnnotator(text_thickness=1)

for frame in frame_generator:
    annotated_frame = frame.copy()

    labels = [
        f"#{tracker_id} {model.class_names[class_id]} {confidence:0.2f}"
        for _, _, confidence, class_id, tracker_id, _
        in detections
    ]

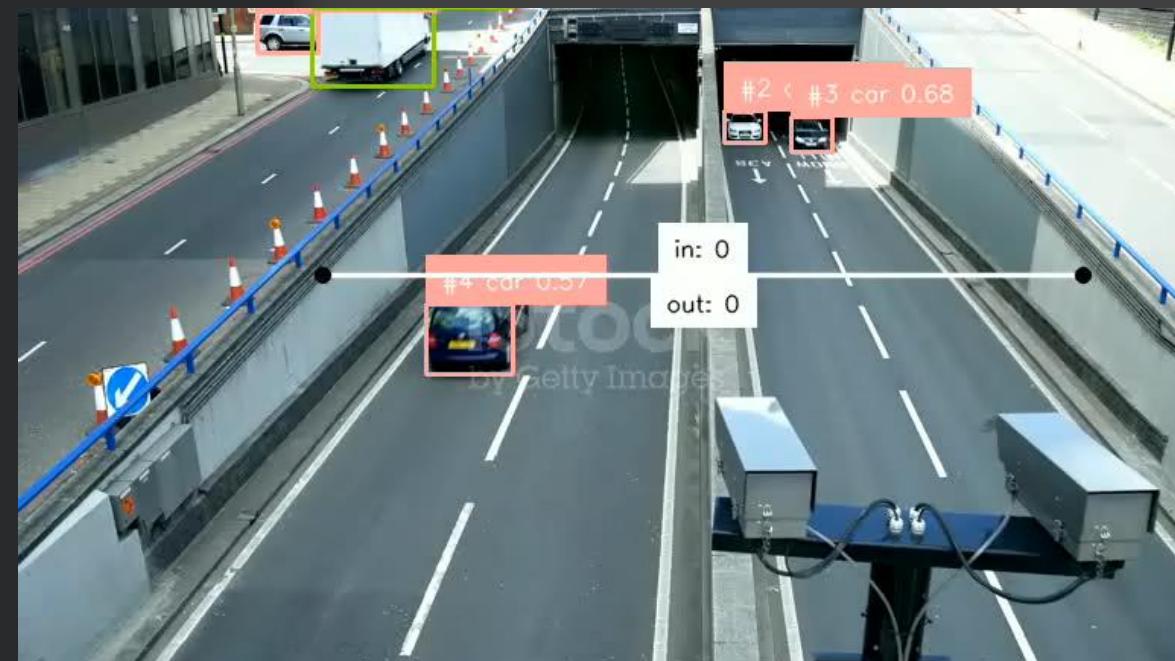
    annotated_frame = box_annotator.annotate(
        annotated_frame,
        detections)

    annotated_frame = trace_annotator.annotate(
        annotated_frame,
        detections)

    annotated_frame = label_annotator.annotate(
        scene=annotated_frame,
        detections=detections,
        labels=labels)

    annotated_frame = line_zone_annotator.annotate(
        annotated_frame,
        line_counter=line_zone)

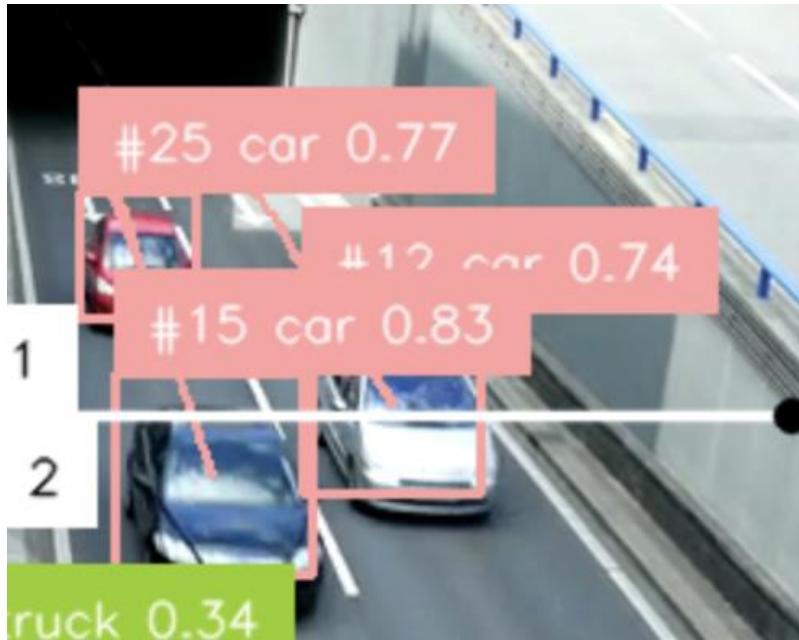
cv2.imshow("", annotated_frame)
if cv2.waitKey(1) & 0xFF == ord("q"):
    break
```



everything YOLO knows

person	elephant	wine glass	dining table
bicycle	bear	cup	toilet
car	zebra	fork	tv
motorcycle	giraffe	knife	laptop
airplane	backpack	spoon	mouse
bus	umbrella	bowl	remote
train	handbag	banana	keyboard
truck	tie	apple	cell phone
boat	suitcase	sandwich	microwave
traffic light	frisbee	orange	oven
fire hydrant	skis	broccoli	toaster
stop sign	snowboard	carrot	sink
parking meter	sports ball	hot dog	refrigerator
bench	kite	pizza	book
bird	baseball bat	donut	clock
cat	baseball glove	cake	vase
dog	skateboard	chair	scissors
horse	surfboard	couch	teddy bear
sheep	tennis racket	potted plant	hair drier
cow	bottle	bed	toothbrush

fine-tuned object detection model



Yellow Taxis
+7 percentage points

Taxi license plates begin with **Y** and end in **C**



For Hire (Includes Uber, Lyft)
-1 pct. pts

For-hire vehicle license plates begin with **T** and end in **C**



Private vehicles
-6 pct. pts

All remaining vehicles were classified as private



Commercial
No change

Includes vans, box trucks and large freight vehicles

The most flexible data labeling platform to fine-tune LLMs, prepare training data or validate AI models.

Quick Start Compare Versions

LAST COMMIT: JANUARY 28, 2025 | LATEST VERSION: 1.15.0

GENAI IMAGES AUDIO TEXT TIME SERIES MULT

Label every data

Discover & Learn About Label Studio Terminology Compare editions

Install & Setup Quick start Installation Set up the database Set up persistent storage Start Label Studio Upgrade Label Studio Secure Label Studio

Create & Manage Projects Create project Configure labeling interface Manage projects Labeling guide

Import & Export Add project storage Import data Export annotations

MACHINE LEARNING Machine learning integration Write your own ML backend Write your own ML Tutorials

Open Source Data Labeling Platform

Quick Start

PIP BREW GIT DOCKER

```
1 # Install the package
# into python virtual environment
2 pip install -U label-studio
3
4 # Launch it!
5 label-studio
```

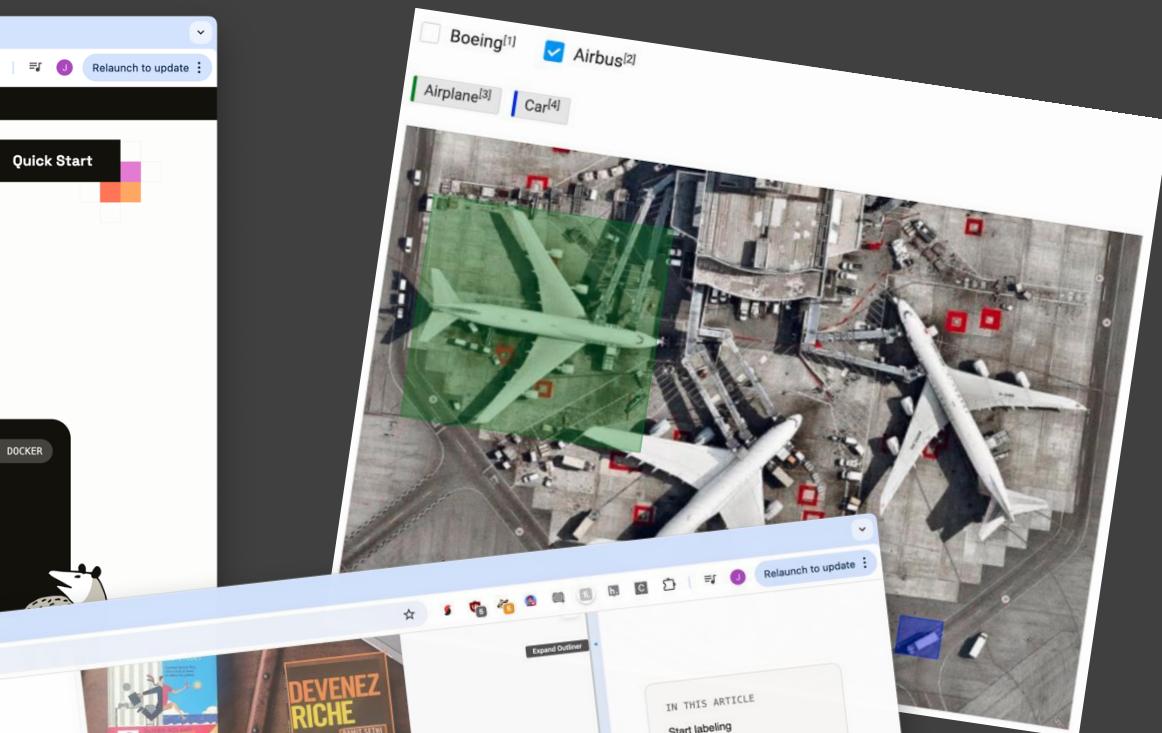


Figure 9: Outliner panel with expand outline icon

The Outliner panel is displayed in expanded view to explore the feature.

IN THIS ARTICLE

- Start labeling
- Choose which tasks to label
- Label a region in the data
- Label overlapping regions
- Select multiple regions
- Duplicate regions
- Change the label
- Delete a labeled region
- Delete an annotation
- Add relations between annotations
- Skipping a task
- Exiting a labeling flow
- Label with collaborators
- Perform ML-assisted labeling with interactive preannotations
- Use keyboard shortcuts
- Image-specific hotkeys

Designed for teams of all sizes

Compare Versions

Annotation History #10000714 Just now

Details

Annotation History #10000714 Just now

Comments Add a comment

Relations (0)

Manual Grouping Ordered by Score

Outline

Figure 10: Outliner panel in expanded view

These panels are detachable, so you can place them wherever you want on the screen, and you can resize them.

fine-tuning

Hugging Face

Models - Hugging Face

huggingface.co/models

Hugging Face

Search models, datasets

Models Datasets Spaces Docs Enterprise

Tasks Libraries Datasets Languages Licenses

Other

Filter Tasks by name

Multimodal

- Audio-Text-to-Text
- Image-Text-to-Text
- Visual Question Answering
- Document Question Answering
- Video-Text-to-Text
- Any-to-Any

Computer Vision

- Depth Estimation
- Image Classification
- Object Detection
- Image Segmentation
- Text-to-Image
- Image-to-Text
- Image-to-Image
- Image-to-Video
- Unconditional Image Generation
- Video Classification
- Text-to-Video
- Zero-Shot Image Classification
- Mask Generation
- Zero-Shot Object Detection
- Text-to-3D
- Image-to-3D
- Image Feature Extraction
- Keypoint Detection

Models 1,345,912

Full-text search

Filter by name

merve (@mervenoyann) / X

25.4K posts

They don't know "Yes we do"
I work at Hugging Face

merve @mervenoyann

open-sourceress at @huggingface 🌐 proud Mediterrenean 🌞 I work on: zero-shot vision & VLMs, large multimodal models, transformers, smol-vision

Paris huggingface.co/merve Joined December 2019

4,593 Following 70.9K Followers

Followed by Ben Burtenshaw, doomslide, and 86 others you follow

Posts Replies Highlights Media

Pinned

merve

https://x.com/mervenoyann/j

Roboflow: Computer

roboflow.com

Sign in Book a demo Get Started

roboflow

Everything you need to build and run computer vision applications.

Request a Demo

Started

Counting Analysis

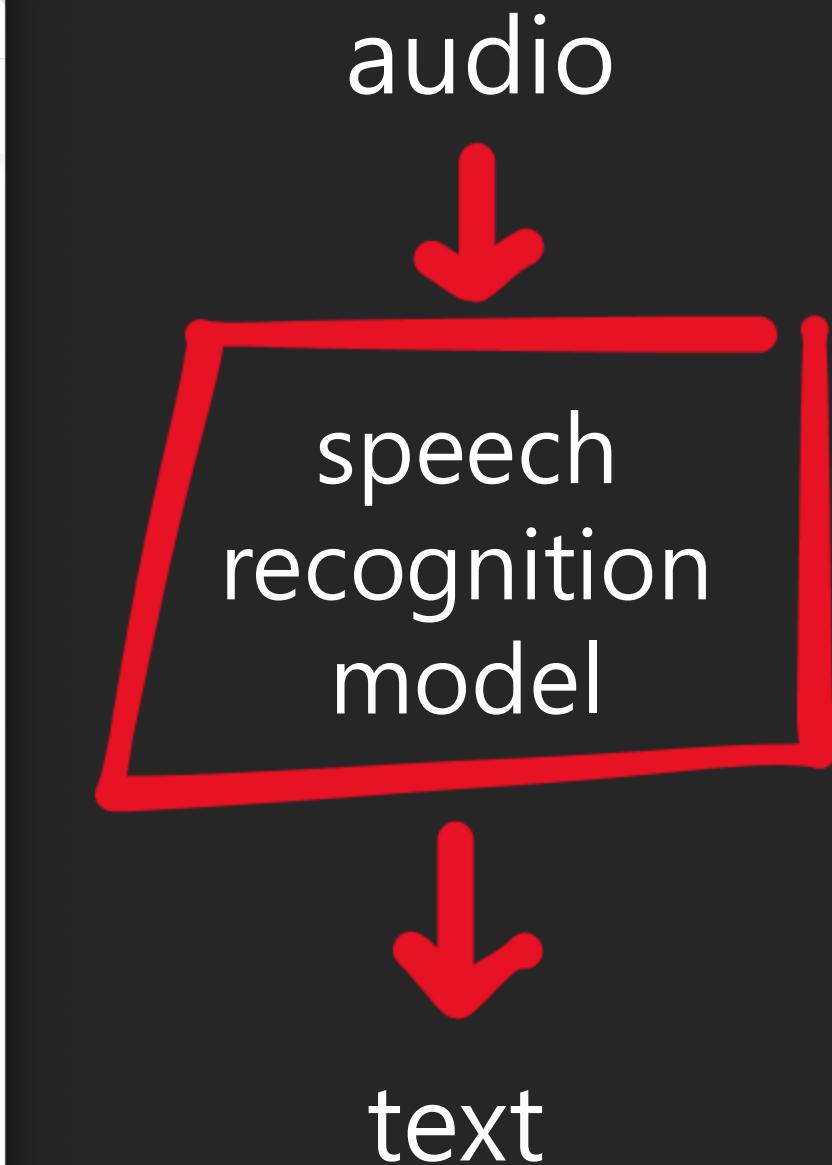
1 trending on notebooks

Misinformation on TikTok: How 'Documented' Examined Hundreds of Videos in Different Languages

Author: Lam Thuy Vo
GRANTEE

SECTIONS MORE

Relaunch to update

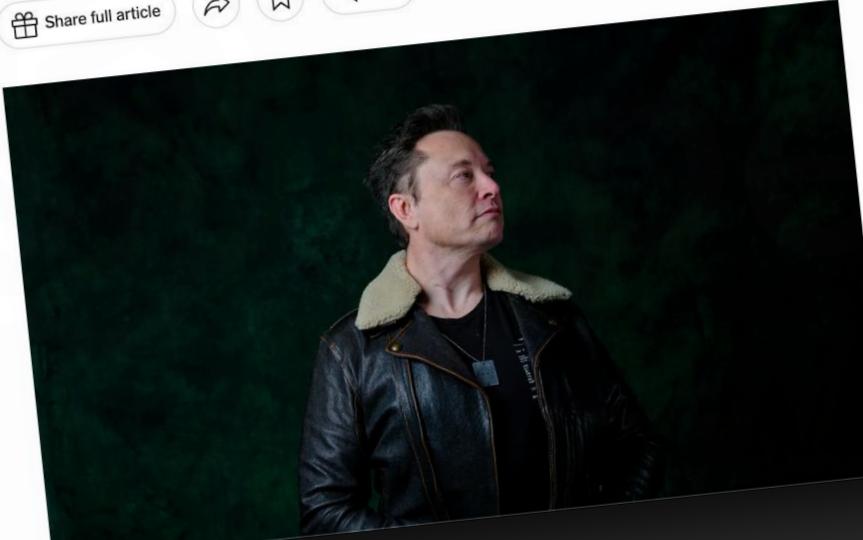


Elon Musk Sues OpenAI and Sam Altman for Violating the Company's Principles

Musk said the prominent A.I. start-up had put profits and commercial interests ahead of seeking to benefit humanity.

Share full article

328



The New York Times

Harvesting Data for A.I. A.I. Data Race, Explained What Is 'Synthetic Data?' Key Figures in the Field A.I. Faces Quiz

ChatGPT

You

How many n's are in mayonnnnaisne?

ChatGPT

The word "mayonnaise" typically has two 'n's in it. However, in the way you've spelled it—"mayonnnnaisne"—there are three 'n's.

AI

ChatGPT...

ChatGPT can make mistakes. Consider checking important information.

The image shows a dual-monitor setup. The left monitor displays the OpenAI Whisper landing page, featuring a large "Introducing Whisper" title, a date of September 21, 2022, and a paragraph about the model's performance. It includes links to "Read paper", "View code", and "View model card". The background is a blurred image of a landscape. The right monitor displays the Hugging Face model library, specifically the "Models" section. It lists several Whisper models by OpenAI, including "openai/whisper-large-v3", "openai/whisper-large-v2", "openai/whisper-tiny", "openai/whisper-base", "openai/whisper-small", "openai/whisper-medium", and "openai/whisper-tiny.en". Each entry shows the model name, type (Automatic Speech Recognition), last update date, size, and the number of stars.

Introducing Whisper | OpenAI

openai.com/index/whisper

Research Products Safety Company

September 21, 2022

Introducing Whisper

We've trained and are open-sourcing a neural net called Whisper that approaches human level robustness and accuracy on English speech recognition.

Read paper ↗ View code ↗ View model card ↗

Models - Hugging Face

huggingface.co/models?sort=trending&sear...

Hugging Face

Search models, datasets, users...

Models 292

openai/whisper

openai/whisper-large-v3
Automatic Speech Recognition • Updated Feb 8 • 2.01M • 2.41k

openai/whisper-large-v2
Automatic Speech Recognition • Updated Feb 29 • 306k • 1.56k

openai/whisper-tiny
Automatic Speech Recognition • Updated Feb 29 • 216k • 180

openai/whisper-base
Automatic Speech Recognition • Updated Feb 29 • 359k • 159

openai/whisper-small
Automatic Speech Recognition • Updated Feb 29 • 368k • 152

openai/whisper-medium
Automatic Speech Recognition • Updated Feb 29 • 188k • 174

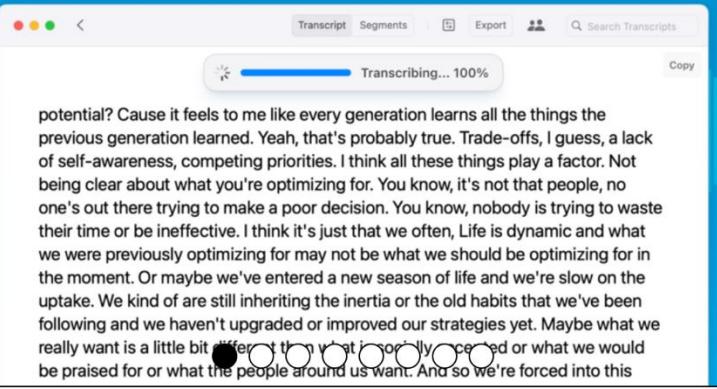
openai/whisper-tiny.en
Automatic Speech Recognition • Updated Jan 22 • 170k • 77

MacWhisper

goodsnooze.gumroad.com/l/mac... ☆

Private, high-quality transcription

100K+ users | Featured by Apple



potential? Cause it feels to me like every generation learns all the things the previous generation learned. Yeah, that's probably true. Trade-offs, I guess, a lack of self-awareness, competing priorities. I think all these things play a factor. Not being clear about what you're optimizing for. You know, it's not that people, no one's out there trying to make a poor decision. You know, nobody is trying to waste their time or be ineffective. I think it's just that we often, Life is dynamic and what we were previously optimizing for may not be what we should be optimizing for in the moment. Or maybe we've entered a new season of life and we're slow on the uptake. We kind of are still inheriting the inertia or the old habits that we've been following and we haven't upgraded or improved our strategies yet. Maybe what we really want is a little bit different than what is currently coded or what we would be praised for or what the people around us want. And so we're forced into this

MacWhisper

Jordi Bruin | ★★★★★ 1268 ratings

Quickly and easily transcribe audio files into text with OpenAI's state-of-the-art transcription technology Whisper. Whether you're recording a meeting, lecture, or other important audio, MacWhisper quickly and accurately transcribes your audio files into text.

€0

I want this!

EasyWhisper | Easily transcript... ☆

easywhisper.io

EASY WHISPER

Advantages Key Features Prices FAQs

Easily transcribe audio files into text

with the advanced transcription technology Whisper

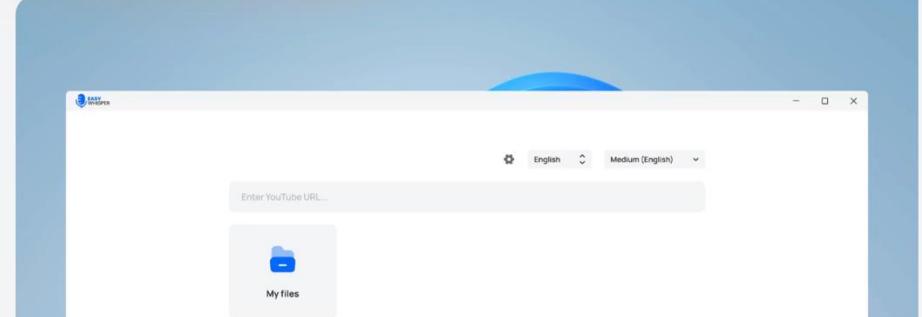
You purchase our app once and for all, **no subscription fee!**

Buy for Windows

Buy for Mac OS

mac OS

Get you trial version here





MacWhisper



Jordi Bruin

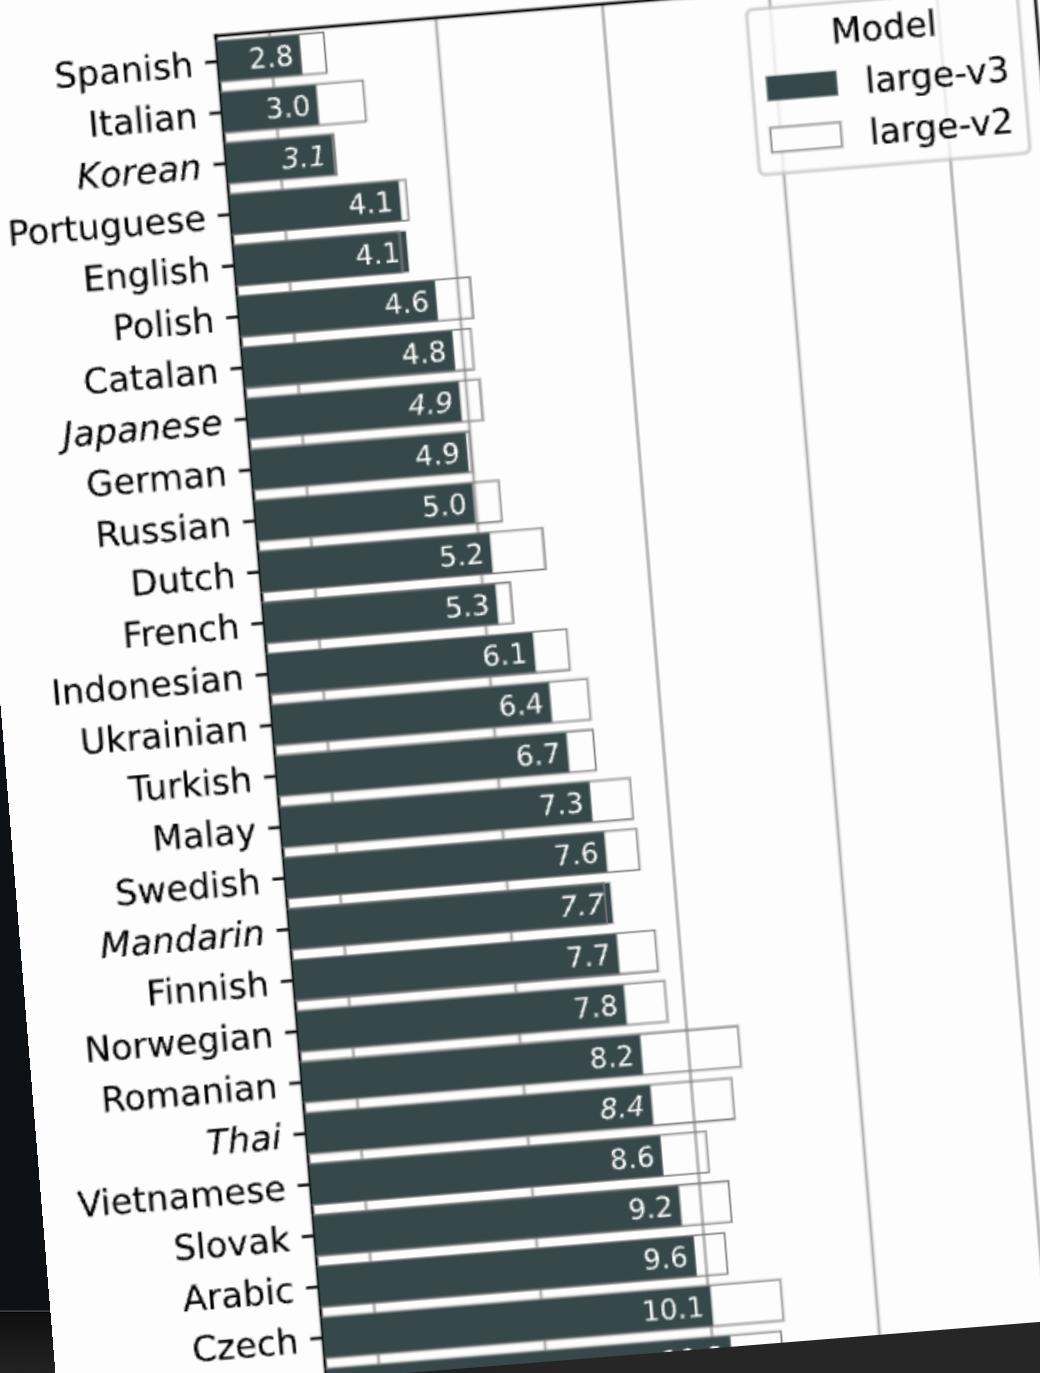
★★★★★ 1268 ratings

Quickly and easily transcribe audio files into text with OpenAI's state-of-the-art transcription technology Whisper. Whether you're recording a meeting, lecture, or other important audio, MacWhisper quickly and accurately transcribes your audio files into text.

Features

- Easily record and transcribe audio files on your Mac
- Just drag and drop audio files to get a high quality transcription
- Supported formats: mp3, wav, m4a, ogg, opus, mov and mp4 videos.
- Adjust whisper settings (beam search / greedy, beam size etc)
- Supports Tiny, Base and Small transcription qualities (and you can compare the transcription quality between all quality levels)





Available models and languages

There are five model sizes, four with English-only versions, offering speed and accuracy tradeoffs. Below are the names of the available models and their approximate memory requirements and inference speed relative to the large model; actual speed may vary depending on many factors including the available hardware.

Size	Parameters	English-only model	Multilingual model	Required VRAM	Relative speed
tiny	39 M	tiny.en	tiny	~1 GB	~32x
base	74 M	base.en	base	~1 GB	~16x
small	244 M	small.en	small	~2 GB	~6x
medium	769 M	medium.en	medium	~5 GB	~2x
large	1550 M	N/A	large	~10 GB	1x

The .en models for English-only applications tend to perform better, especially for the tiny.en and base.en models. We observed that the difference becomes less significant for the small.en and medium.en models.

Whisper's performance varies widely depending on the language. The figure below shows a performance breakdown of large-v3 and large-v2 models by language, using WERs (word error rates) or CER (character error rates, shown in *italic*) evaluated on the Common Voice 15 and Fleurs datasets. Additional WER/CER metrics corresponding to the other models and datasets can be found in Appendix D.1, D.2, and D.4 of [the paper](#), as well as the BLEU (Bilingual Evaluation Understudy) scores for translation in Appendix D.3.

m-bain/whisperX: WhisperX · GitHub

github.com/m-bain/whisperX

WhisperX

Stars 9.1k Issues 398 open License BSD-4-Clause Arxiv 2303.00747 Tweet

Input audio → Cut & Merge → Batch → Whisper → Phoneme Model → Forced Alignment → Transcription word-level timestamps

Voice Activity Detection → Cut & Merge → Pad to 30s

This repository provides fast automatic speech recognition (70x realtime with large-v2) with word-level timestamps and speaker diarization.

- ⚡ Batched inference for 70x realtime transcription using whisper large-v2
- ↗ faster-whisper backend, requires <8GB gpu memory for large-v2 with beam_size=5
- ⌚ Accurate word-level timestamps using wav2vec2 alignment
- 喤 Multispeaker ASR using speaker diarization from [pyannote-audio](#) (speaker ID labels)
- 👉 VAD preprocessing, reduces hallucination & batching with no WER degradation

Whisper is an ASR model [developed by OpenAI](#), trained on a large dataset of diverse audio. It does produce highly accurate transcriptions, the corresponding timestamps are at the utterance level, not per word, and can be inaccurate by several seconds. OpenAI's whisper does not natively support batching.

Phoneme-Based ASR A suite of models finetuned to recognise the smallest unit of speech distinguishing one word from another, e.g. the element p in "tap". A popular example model is [wav2vec2.0](#).

Forced Alignment refers to the process by which orthographic transcriptions are aligned to audio recordings to automatically generate phone-level segmentation.

+ 6 releases

Sponsor this project

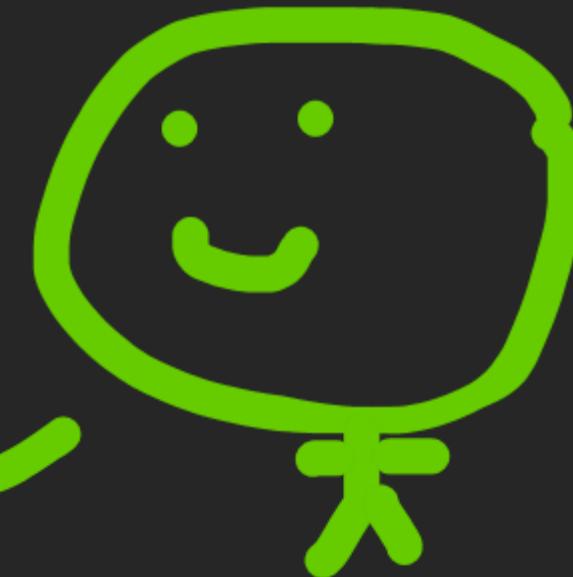
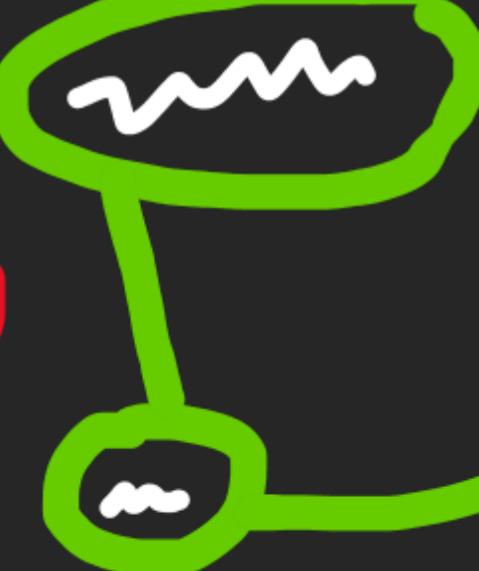
<https://www.buymeacoffee.com/maxh...>

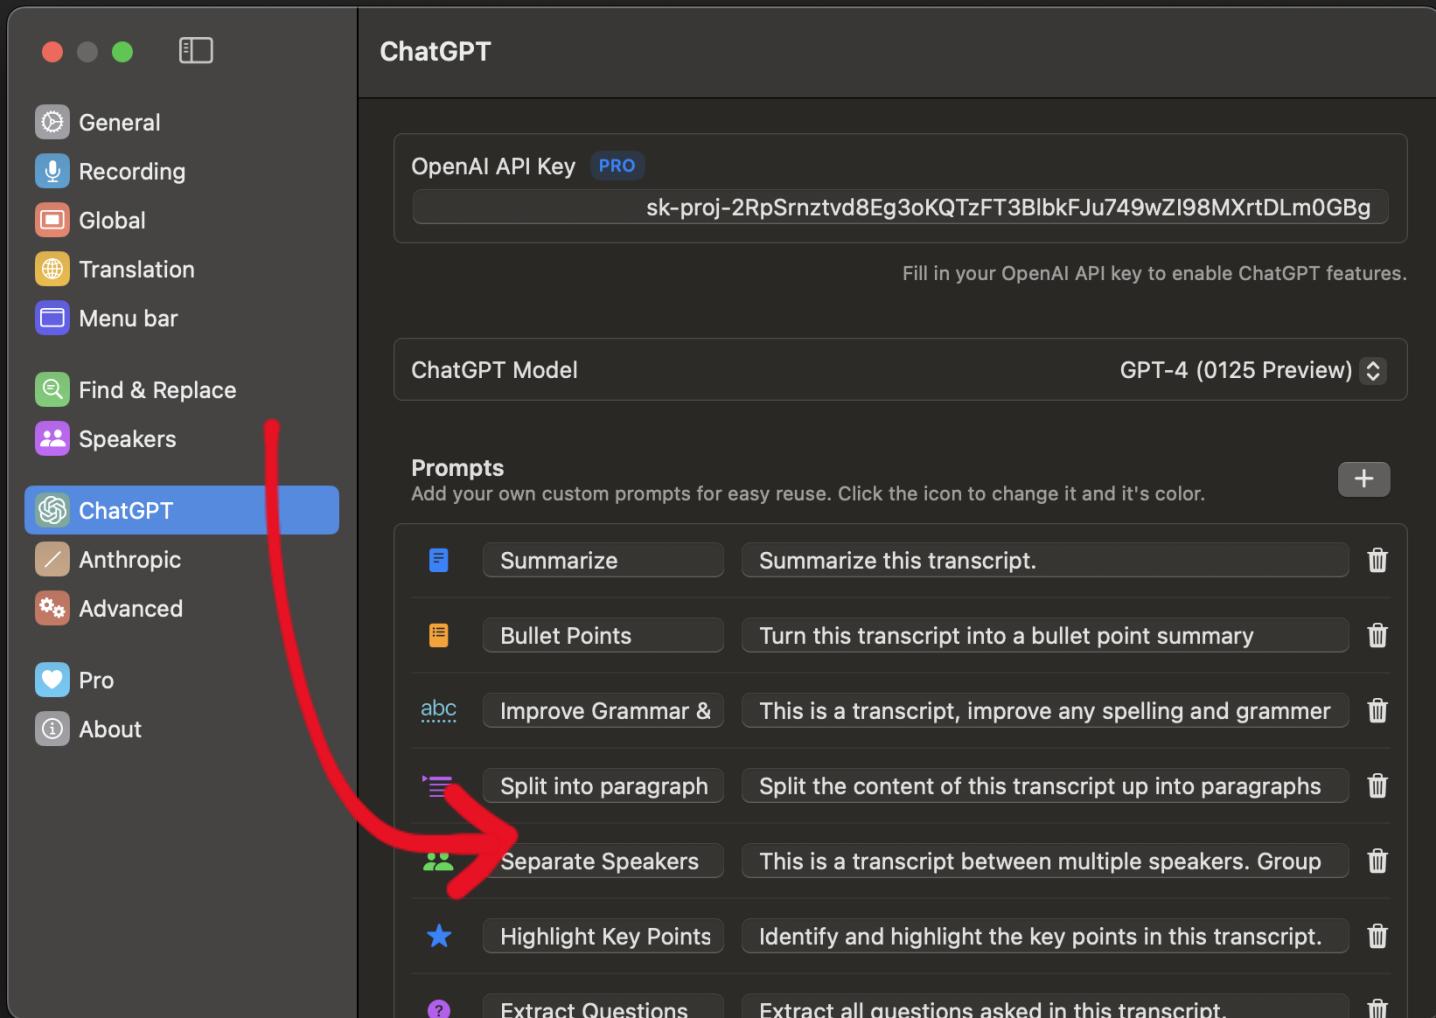
```
import whisperx
import torch

device = "cuda" if torch.cuda.is_available() else "cpu"
compute_type = "int8"

model = whisperx.load_model("base",
                            device,
                            compute_type=compute_type)
audio = whisperx.load_audio("audio.mp3")
result = model.transcribe(audio, batch_size=4)
```

Diarization is hard





This is a transcript between multiple speakers. Group sentences said by the same person together and separate each section with a new line.

AP Researchers say AI transcript X + apnews.com/article/ai-artificial-intelligence-health-business... Relaunch to update

WORLD U.S. POLITICS SPORTS ENTERTAINMENT BUSINESS SCIENCE FACT CHECK ODDITIES BE WELL ...

Live: DC plane crash Gaza ceasefire Guantanamo Bay Kumbh Mela stampede RFK J AP SETS THE STANDARD FOR POLITICAL REPORTING. SUPPORT INDEPENDENT, FACT-BASED JOURNALISM. DONATE

TECHNOLOGY

Researchers say an AI-powered transcription tool used in hospitals invents things no one ever said

```
text": "Thank you."
}
# Ground truth: "a cat to cat cat"
audio_file= open("cat.wav", "rb")
print(openai.Audio.transcribe("whisper-1", audio_file))
8.9s
{
    "text": "The guy in the black coat, there was one
}
# Ground truth: "well if I was gonna make a sandwich with ou
# I really I prefer a really good bakery that has"
audio_file= open("sandwich.wav", "rb")
print(openai.Audio.transcribe("whisper-1", audio_file))
[10]
...
{
    "text": "Well, if I was going to make a sandwich out of peanuts and some kind of fruit, I would go to
prefer a really good bakery that has a really good sandwich."
}
# Ground truth: "as he s going out to off to the second picture she's given a little I think say"
audio_file= open("picture.wav", "rb")
print(openai.Audio.transcribe("whisper-1", audio_file))
Jupyter Server: Local Cell 6 of 9
```

1 of 6 | Assistant professor of information science Allison Koencke, an author of a recent study that found hallucinations in a speech-to-text transcription tool, works in her office at Cornell University in Ithaca, N.Y., Friday, Feb. 2, 2024. The text preceded by "#Ground truth" Read More

In an example they uncovered, a speaker said, “He, the boy, was going to, I’m not sure exactly, take the umbrella.”

But the transcription software added: “He took a big piece of a cross, a teeny, small piece … I’m sure he didn’t have a terror knife so he killed a number of people.”

Trump White House rescinds memo freezing federal money after widespread confusion

Robert F. Kennedy Jr. struggles to answer questions on Medicare and Medicaid at confirmation hearing

What is the Laken Riley Act? A look at the first bill

	whisperx.large-v3	whispercpp.large-v3	insanely-fast-whisper.large-v3	macwhisper.large	youtube	deepgram	whisperx.base	insanely-fast-whisper.base	otter_ai	insanely-fast-whisper.tiny
whisperx.large-v3	0.00	0.14	0.15	0.17	0.26	0.28	0.30	0.37	0.47	0.51
whispercpp.large-v3	0.14	0.00	0.15	0.16	0.23	0.23	0.28	0.34	0.44	0.48
insanely-fast-whisper.large-v3	0.15	0.15	0.00	0.18	0.28	0.29	0.32	0.36	0.46	0.50
macwhisper.large	0.17	0.17	0.19	0.00	0.28	0.29	0.31	0.37	0.47	0.52
deepgram	0.26	0.23	0.29	0.27	0.28	0.00	0.32	0.37	0.44	0.48
youtube	0.26	0.24	0.29	0.28	0.00	0.29	0.31	0.36	0.42	0.50
whisperx.base	0.31	0.30	0.34	0.31	0.32	0.35	0.00	0.28	0.44	0.51
insanely-fast-whisper.base	0.36	0.35	0.37	0.36	0.36	0.39	0.27	0.00	0.47	0.48
insanely-fast-whisper.tiny	0.48	0.49	0.50	0.49	0.48	0.48	0.47	0.46	0.52	0.00
otter_ai	0.54	0.54	0.56	0.54	0.49	0.54	0.50	0.55	0.00	0.63

You can use the WER calculation from the machine recognition results to evaluate the quality of the model you're using with your app, tool, or product. A WER of 5-10% is considered to be good quality and is ready to use. A WER of 20% is acceptable, but you might want to consider more training. A WER of 30% or more signals poor quality and requires customization and training.

-Microsoft

click me

Models - Hugging Face Relaunch to update

huggingface.co/models

Search models, datasets, users...

Hugging Face

Models 1,353,946 Filter by name

Full-text search Sort: Trending

Models

Datasets Spaces Posts Docs Enterprise Pricing

Tasks Libraries Datasets Languages Licenses Other

Filter Tasks by name

Multimodal

- Audio-Text-to-Text
- Image-Text-to-Text
- Visual Question Answering
- Document Question Answering
- Video-Text-to-Text
- Any-to-Any

Computer Vision

- Depth Estimation
- Object Detection
- Image Classification
- Image Segmentation
- Text-to-Image
- Image-to-Text
- Image-to-Image
- Image-to-Video
- Unconditional Image Generation
- Video Classification
- Text-to-Video
- Zero-Shot Image Classification
- Mask Generation
- Zero-Shot Object Detection
- Text-to-3D
- Image-to-3D
- Image Feature Extraction
- Keypoint Detection

Natural Language Processing

- Text Classification
- Token Classification
- Table Question Answering
- Question Answering
- Zero-Shot Classification
- Translation
- Summarization
- Feature Extraction
- Text Generation
- Text2Text Generation
- Fill-Mask
- Sentence Similarity

https://huggingface.co/deepseek-ai/DeepSeek-V3

deepseek-ai/DeepSeek-R1

Text Generation Updated 4 days ago ↓ 498k ⚡ 5.14k

deepseek-ai/DeepSeek-V3

Text Generation Updated 6 days ago ↓ 409k ⚡ 2.82k

deepseek-ai/DeepSeek-R1-Distill-Qwen-32B

Text Generation Updated 4 days ago ↓ 194k ⚡ 735

hexgrad/Kokoro-82M

Text-to-Speech Updated about 14 hours ago ↓ 55.4k ⚡ 2.57k

deepseek-ai/Janus-Pro-1B

Any-to-Any Updated 3 days ago ↓ 17.5k ⚡ 282

deepseek-ai/DeepSeek-R1-Zero

Text Generation Updated 4 days ago ↓ 17.1k ⚡ 617

Qwen/Qwen2.5-VL-7B-Instruct

Image-Text-to-Text Updated 3 days ago ↓ 36.2k ⚡ 201

deepseek-ai/DeepSeek-R1-Distill-Llama-8B

Text Generation Updated 4 days ago ↓ 153k ⚡ 344

Qwen/Qwen2.5-14B-Instruct-1M

Text Generation Updated 1 day ago ↓ 3.85k ⚡ 181

deepseek-ai/DeepSeek-V3-Base

Updated 6 days ago ↓ 23.4k ⚡ 1.46k

deepseek-ai/Janus-Pro-7B

Any-to-Any Updated 3 days ago ↓ 79.3k ⚡ 2.03k

tencent/Hunyuan3D-2

Image-to-3D Updated 6 days ago ↓ 24.2k ⚡ 640

HKUSTAudio/Lasa-3B

Text-to-Speech Updated 4 days ago ↓ 4.6k ⚡ 381

unsloth/DeepSeek-R1-GGUF

Text Generation Updated 1 day ago ↓ 209k ⚡ 341

deepseek-ai/DeepSeek-R1-Distill-Qwen-1.5B

Text Generation Updated 4 days ago ↓ 22.5k ⚡ 562

m-a-p/YuE-s1-7B-anneal-en-cot

Text Generation Updated 11 minutes ago ↓ 5.27k ⚡ 242

cyberagent/DeepSeek-R1-Distill-Qwen-32B-Japanese

Text Generation Updated 3 days ago ↓ 4.77k ⚡ 193

deepseek-ai/DeepSeek-R1-Distill-Llama-70B

Text Generation Updated 4 days ago ↓ 92.7k ⚡ 376

black-forest-labs/FLUX.1-dev

Text-to-Image Updated Aug 16, 2024 ↓ 1.54M ⚡ 8.35k

Qwen/Qwen2.5-VL-72B-Instruct

Image-Text-to-Text Updated 3 days ago ↓ 9.47k ⚡ 161

models

"AI"



youtube.com/playlist?list=PLewNEVDy7gq1_GPUaL0Q31QsiHP5ncAQ

YouTube BE

Search

Home

Shorts

Subscriptions

You >

Your channel

History

Playlists

Your videos

Watch later

Liked videos

Your clips

Dearest friend, I daresay I have not partaken of food in ages. I'm positively famished

Practical AI for Investigative Journalism

Jonathan Soma

Playlist • Public • 6 videos • 669 views

A six session series held in April 2024 about real-life use cases for journalism in (mostly investigative) jour...more

▶ Play all

Sort

Dearest friend, I daresay I have not partaken of food in ages. I'm positively famished 2:09:33

Sorting documents (Practical AI for Investigative Journalism, Session 1)

Jonathan Soma • 1.1K views • Streamed 2 months ago

Large language models don't understand facts or concepts, they only know statistical probability 2:14:32

Structured, validated data from LLMs (Practical AI for Investigative Journalism, Session 2)

Jonathan Soma • 743 views • Streamed 2 months ago

Evaluating Verifiability in Generative Search Engines 1:45:26

Why generative AI is a dead end for responsible journalism (Practical AI for Journalism, Session 3)

Jonathan Soma • 581 views • Streamed 2 months ago

AI, Hugging Face and non-chatbot models (Practical AI for Journalism, Session 4)

Jonathan Soma • 836 views • Streamed 2 months ago

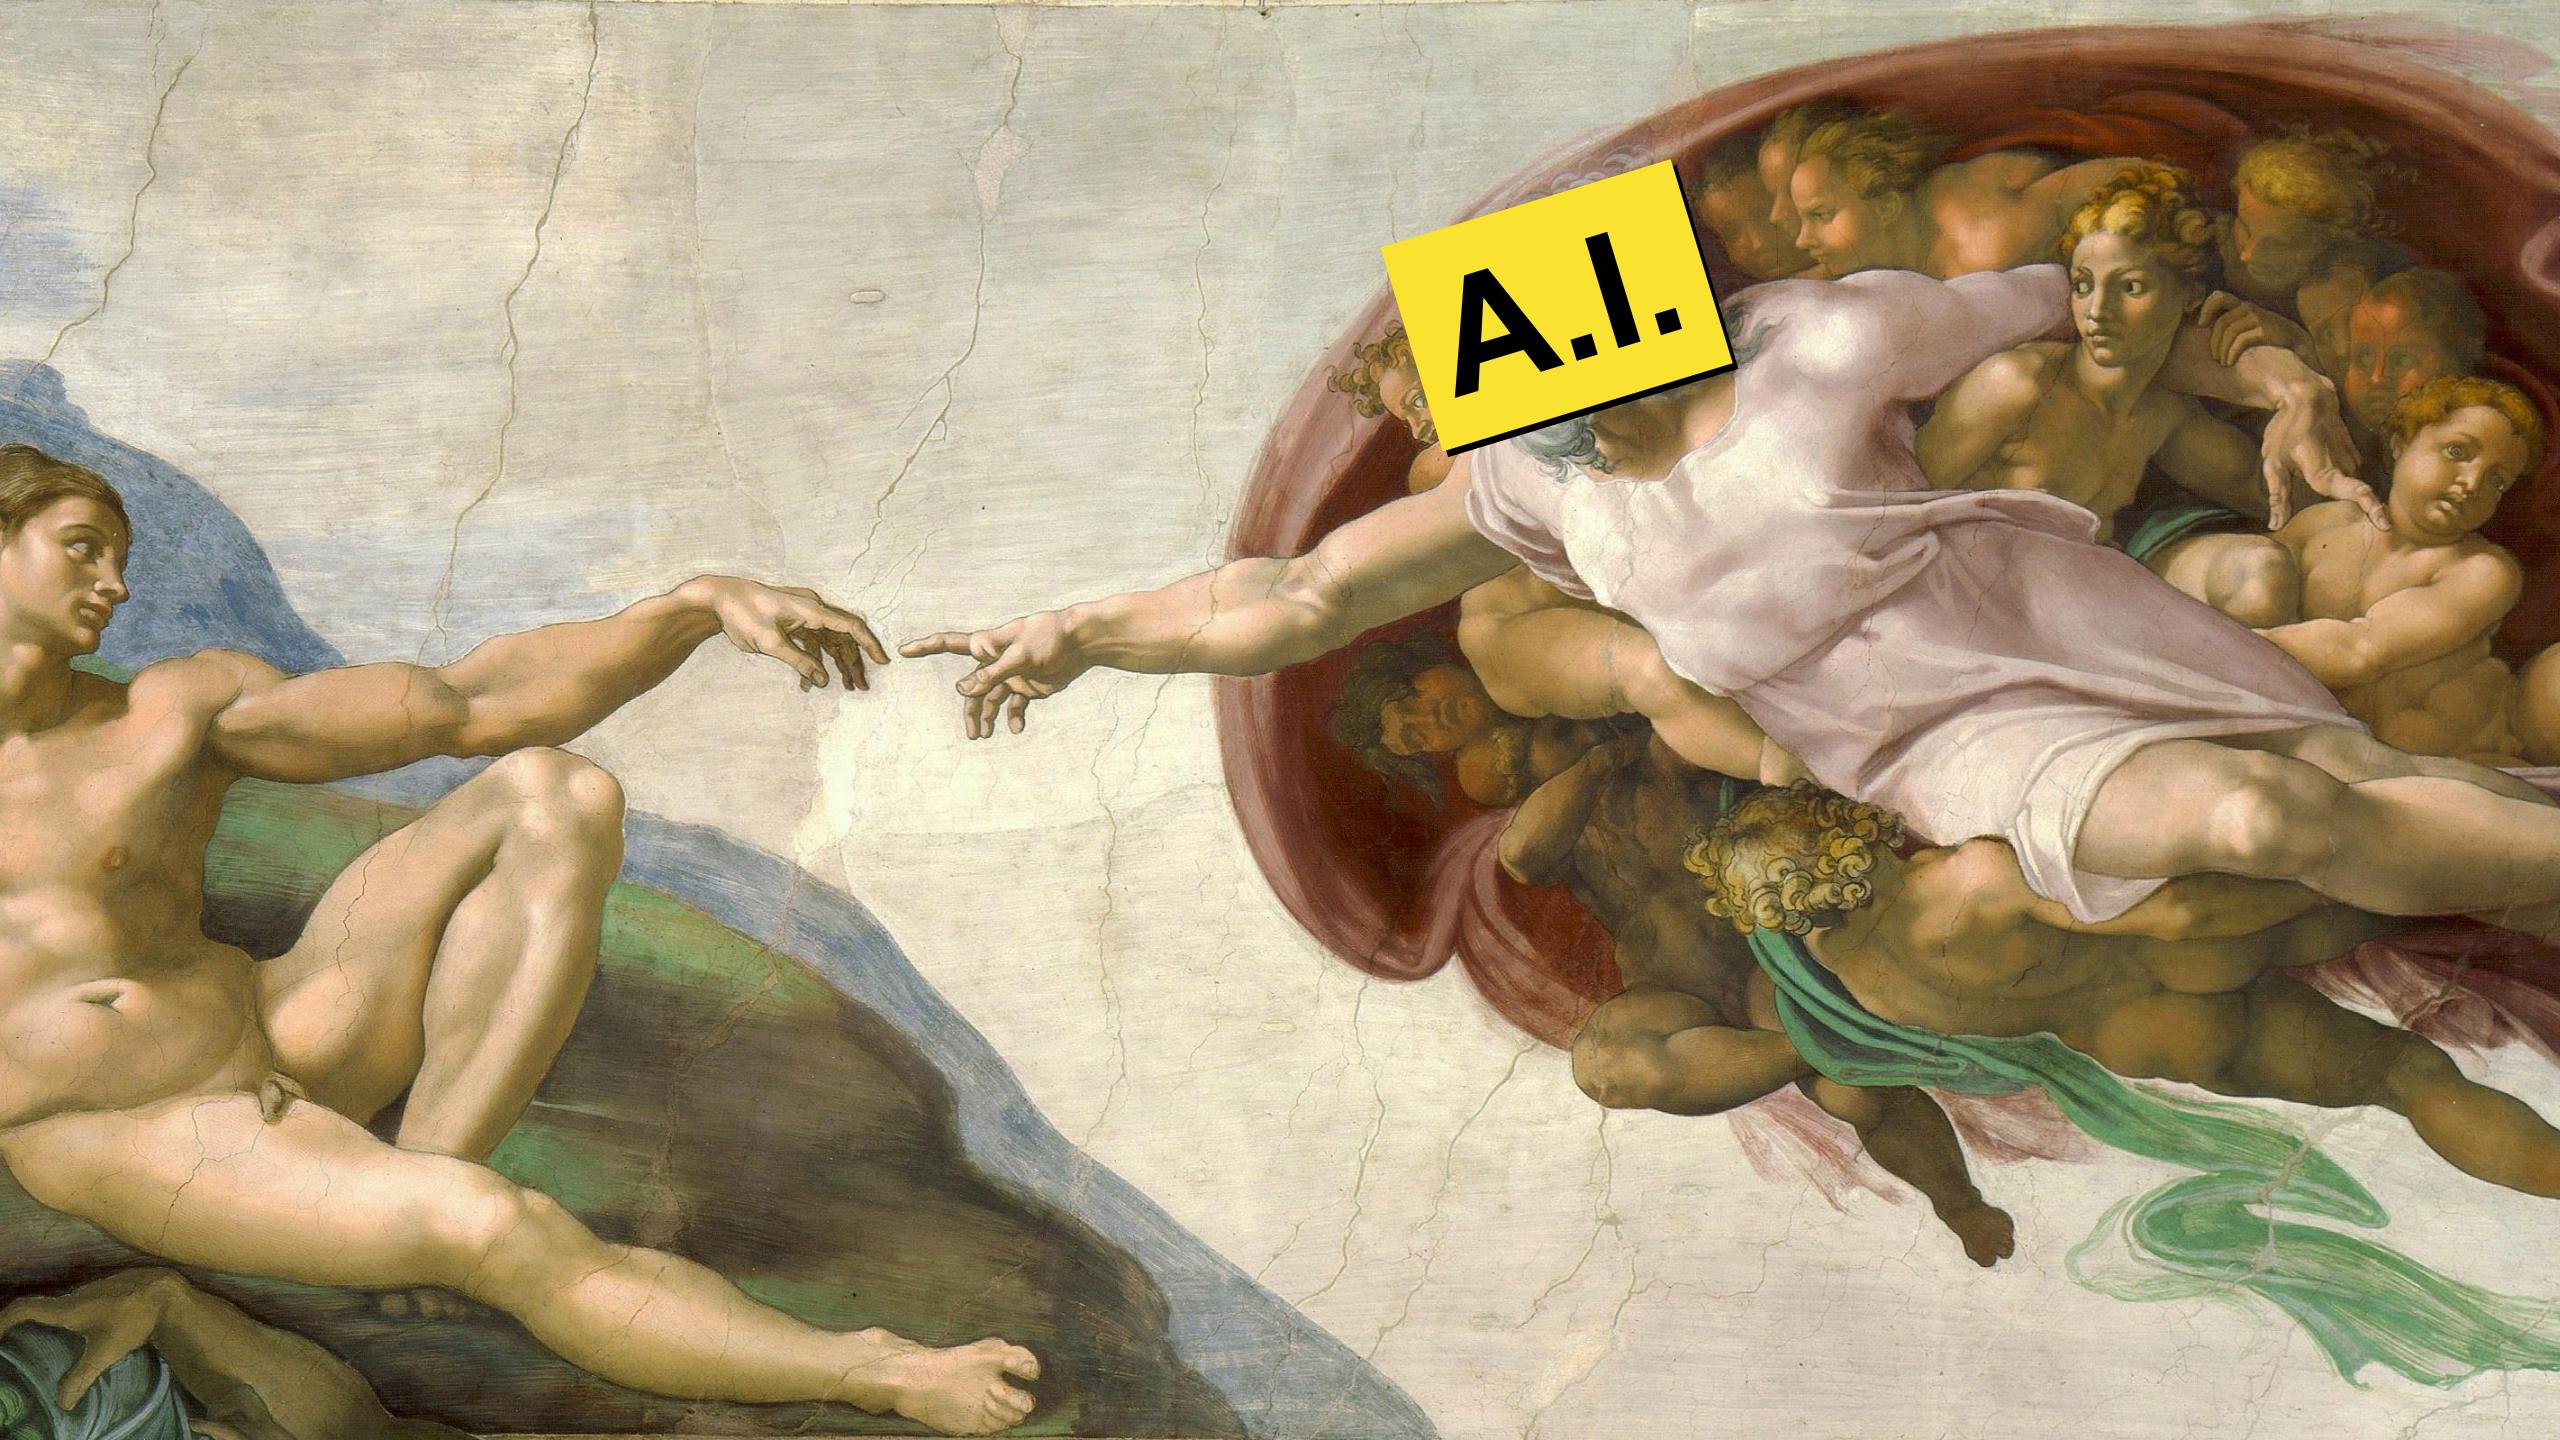
Context window 2:01:02

Local models/private AI (Practical AI for Investigative Journalism, Session 5)

Jonathan Soma • 451 views • Streamed 2 months ago

Transcription and audio models (Practical AI for Investigative Journalism, Session 6)

Jonathan Soma • 332 views • Streamed 2 months ago



A.I.

A whole bushel of AI things

<https://bit.ly/soma-econ>

Jonathan Soma

Knight Chair in Data Journalism

Columbia University Graduate School of Journalism

jonathansoma.com :: @dangerscarf :: js4571@columbia.edu